



Trinity College Dublin
Coláiste na Tríonóide, Baile Átha Cliath
The University of Dublin

Assignment 3

CS7IS2 Artificial Intelligence

Ujjayant Kadian (22330954)

April 5, 2025

Introduction

This report explores the implementation and evaluation of artificial intelligence agents in two well-known board games: Tic Tac Toe and Connect 4. The main focus is on comparing the performance of traditional search-based algorithms, such as Minimax (with and without alpha-beta pruning), against reinforcement learning methods, specifically Q-Learning.

Both games were developed from scratch using modular, object-oriented Python code, which allowed for a clear separation between the core game logic and the graphical interface. The integration of AI agents was followed by systematic experiments under various conditions, including matches against semi-intelligent default opponents and direct competitions between the AI agents.

The evaluation highlights significant trade-offs between learning-based and deterministic approaches, considering factors such as strategy depth, scalability, training overhead, and real-time performance. The findings provide insights into the suitability of each algorithm type for different problem scales and complexity levels, while also addressing specific implementation nuances and performance bottlenecks.

Game Implementations

General Setup

The games, Tic Tac Toe and Connect 4, were implemented in Python, leveraging `numpy` for efficient board state manipulation and `pygame` for rendering interactive graphical interfaces. Both games were implemented entirely from scratch rather than adapting open-source code, enabling complete control over game logic and integration with AI agents. This approach was justified by the necessity to clearly separate and customize the core mechanics, facilitating a clean integration with different algorithms (Minimax, Q-Learning, Semi-Intelligent).

Tic Tac Toe

Board Representation and Rules

The Tic Tac Toe board is represented as a 3x3 `numpy` array initialized with zeros. Player 1 (X) and Player 2 (O) are represented numerically as 1 and 2, respectively. The game adheres strictly to classic Tic Tac Toe rules, where players alternate placing their symbols aiming to align three horizontally, vertically, or diagonally.

Move Validation and Action Space

Valid moves are identified by inspecting unoccupied cells in the `numpy` array. The action space consists of tuples representing cell coordinates (row, col) for each empty space on the board.

Game Loop and Turn Handling

The game loop handles alternating player turns, updates the board state upon valid moves, and checks game termination conditions. It uses clearly defined methods such as `make_move(move)` for updating game states, `_check_game_over()` to verify end conditions, and `get_winner()` to identify the winner.

Connect 4

Board Representation and Gravity Mechanic

Connect 4 employs a 6x7 numpy array to represent the board, initialized with zeros. Each player's pieces, marked numerically as 1 or 2, "drop" to the lowest available space within a selected column, mimicking the gravity-based mechanics of the traditional Connect 4 game.

Move Validation and Win/Draw Logic

Moves are validated by checking column availability. The method `get_legal_moves()` identifies columns with open slots. The `_check_game_over()` method assesses the board after each move for vertical, horizontal, or diagonal connect-four conditions, or a draw scenario if no moves remain.

Game Loop Structure

The Connect 4 game loop sequentially manages player turns, processes moves using `make_move(col)`, and updates the game state. The loop includes logic for animation effects via the `_start_animation(col)` and `_update_animation()` methods, enhancing visual feedback during gameplay.

Reusability and Modularity

Separation of Game Logic and UI

A deliberate design decision was to separate game logic from the user interface (UI) by encapsulating core mechanics within dedicated classes (`TicTacToe` and `Connect4`) independent of graphical handling (`TicTacToeUI` and `Connect4UI`). This modular structure enables easy maintenance and allows the core logic to be reused or adapted for different interfaces or integrations without affecting the gameplay integrity.

Unified Player Types and Game Modes

Both games share common enumerations for player types (`HUMAN`, `AI`, `SEMI_INTELLIGENT`) and game modes (`HUMAN_VS_HUMAN`, `HUMAN_VS_AI`, `HUMAN_VS_SEMI`, `AI_VS_SEMI`, `AI_VS_AI`). The game modes define the interaction dynamics: `HUMAN_VS_HUMAN` allows two players to compete against each other, `HUMAN_VS_AI` pits a human player against an AI opponent, `HUMAN_VS_SEMI` features a human player against a semi-intelligent AI, `AI_VS_SEMI` involves a semi-intelligent AI competing against a fully intelligent AI, and `AI_VS_AI` allows two AI agents to play against each other. This design simplifies configuration management and provides consistent player behavior across both games, effectively supporting various experimental setups involving Minimax, Q-Learning, and Semi-Intelligent agents.

Interface and Interaction

The game classes provide clearly defined methods (`get_state()`, `get_legal_moves()`, `make_move()`, `is_game_over()`, `get_winner()`) that standardize interaction with AI agents. Agents can integrate seamlessly by using these methods to observe the game state, validate and execute moves, and assess game termination. The effectiveness of this design choice lies in its simplicity and clarity, which significantly ease the development and testing of different AI strategies while ensuring straightforward compatibility across both game environments.

Overall, these design decisions and implementations effectively supported modularity, facilitated diverse integration, and allowed focused analysis on algorithmic performance, crucial for the comparative assessment of Minimax and Q-Learning strategies explored in subsequent sections.

Algorithm Implementations

Overview

Two distinct algorithms were implemented and systematically compared: **Minimax** (with optional alpha-beta pruning) and **Q-Learning**. Minimax is a deterministic, exhaustive search algorithm suitable for finite, turn-based adversarial games, systematically exploring game states to make optimal decisions. In contrast, Q-Learning is a reinforcement learning approach, where optimal strategies emerge from repeated gameplay and feedback-based learning. This fundamental contrast: between exhaustive search and incremental learning enables meaningful comparative insights.

Minimax Algorithm

The core idea behind Minimax is to minimize the possible loss for a worst-case scenario. When it's the player's turn, they will choose the move that maximizes their minimum gain (hence "minimax"). Conversely, when it's the opponent's turn, they will choose the move that minimizes the player's maximum gain.

In this context, the game tree represents all possible moves in the game, where each node corresponds to a game state. The root node represents the current state of the game, and each child node represents a possible move leading to a new game state. The depth of the tree refers to how many moves ahead the algorithm evaluates; a greater depth allows for a more thorough analysis of potential future game states.

Pruning, specifically alpha-beta pruning, is an optimization technique for the Minimax algorithm that reduces the number of nodes evaluated in the game tree. It eliminates branches that cannot possibly influence the final decision, allowing the algorithm to run more efficiently by skipping unnecessary calculations.

Tic Tac Toe

Minimax for Tic Tac Toe was implemented in both alpha-beta pruning and non-pruning versions to facilitate experimentation. Given Tic Tac Toe's manageable state space, a complete exploration was feasible with a maximum depth of nine moves.

Evaluation Heuristic:

- **Terminal States:** +100 for wins, -100 for losses, 0 for draws.
- **Intermediate States:** Heuristic scores were assigned based on potential winning alignments, particularly rewarding configurations such as "two-in-a-row" opportunities.

Why This Heuristic Works: Tic Tac Toe's simplicity allows precise evaluations, where each potential alignment strongly influences the outcome. By emphasizing near-complete alignments, the Minimax can decisively exploit immediate threats or opportunities.

Connect 4

Due to Connect 4's complexity, the Minimax algorithm was implemented with adjustable maximum depths to maintain computational feasibility, alongside optional alpha-beta pruning to significantly reduce computational overhead.

Evaluation Heuristic:

- **Window-based evaluation:** Assesses four-cell windows horizontally, vertically, and diagonally, assigning points proportional to the number of pieces a player controls within each window.
- **Positional weighting:** Strongly favors central columns, aligning with known Connect 4 strategies for maximizing potential winning combinations.

Why This Heuristic Works: Connect 4 strategies commonly revolve around central control and creating multiple simultaneous threats. Evaluating windows effectively captures immediate threats and strategic potential, ensuring the Minimax selects optimal strategic moves (like in the case of tic tac toe).

Q-Learning Algorithm

As explained earlier, Q-Learning is a reinforcement learning algorithm that enables agents to learn optimal actions through trial and error by interacting with their environment. The algorithm utilizes a Q-table, which is a data structure that maps each possible game state to a set of action values, representing the expected utility of taking each action from that state. The Q-table effectively stores the rewards associated with each action taken in a given state, allowing agents to make informed decisions based on previously learned

outcomes. Agents were thoroughly trained across thousands of episodes before actual evaluation against the default opponents or the minimax agent, as untrained Q-Learning agents default to random behavior due to uninitialized Q-tables.

Hyperparameters

Hyperparameters are crucial settings that influence the learning process of Q-Learning agents. They help balance the trade-offs between learning speed and the quality of the learned policy.

Learning Rate (α): This parameter determines how much new information overrides old information. A higher learning rate means the agent learns quickly from new experiences, while a lower rate allows for more gradual learning.

- Tic Tac Toe: 0.3, Connect 4: 0.2.

Discount Factor (γ): This factor represents the importance of future rewards. A value closer to 1 emphasizes future rewards, encouraging the agent to consider long-term benefits, while a value closer to 0 focuses on immediate rewards.

- Tic Tac Toe: 0.9, Connect 4: 0.95.

Epsilon (ϵ) (Exploration Rate): This parameter controls the exploration-exploitation trade-off. Exploration involves trying new actions to discover their effects, while exploitation involves choosing the best-known actions based on current knowledge. Starting at 0.3 and decaying to 0.01 allows the agent to initially explore a variety of actions and gradually focus on exploiting the best strategies as it learns.

Reward Structure

Designed to incentivize efficient and effective gameplay:

- **Win:** +1.0
- **Loss:** -1.0
- **Draw:**
 - Tic Tac Toe: +0.2 (as draws are relatively common and better than losses)
 - Connect 4: 0.0
- **Per-move penalty:**
 - Tic Tac Toe: -0.05 (to encourage rapid wins)
 - Connect 4: -0.01 (to slightly favor quicker resolutions)

Game-Specific Heuristics

Tic Tac Toe:

- **Symmetry Exploitation:** The leveraged board symmetries, greatly accelerating learning by applying knowledge gained from one state to equivalent symmetric states. This significantly reduced the Q-table size and improved learning efficiency. (Similar to the minimax agent)

Connect 4:

- **Immediate Threat Detection:** Identifies potential winning moves and opponent threats, allowing the agent to immediately capitalize or defend.
- **Double Threat Detection:** Prioritizes moves that simultaneously create multiple threats, increasing the strategic complexity of moves.
- **Center Control:** Explicitly rewards occupying the central column, a critical strategic advantage in Connect 4.

These heuristics are essentially the same as for Minimax agents. They dramatically enhanced Q-learning performance, aligning learned strategies closely with human-like strategic reasoning.

Default Opponent Agents

Semi-intelligent default opponents served as challenging baselines, superior to random agents due to their incorporation of fundamental game-specific strategies.

Tic Tac Toe Default Opponent:

- Immediately seizes winning opportunities or blocks imminent threats.
- Prefers central positions and corners to edges, reflecting basic strategic principles.

Connect 4 Default Opponent:

- Checks for immediate winning or blocking moves first.
- Prioritizes the central columns, recognizing their strategic value.
- Resorts to random moves only when no strategic advantage or threat is evident.

These default opponents provided more realistic benchmarks, enhancing the experimental validity of evaluating Minimax and Q-Learning agents.

Common Interfaces and Interaction

All algorithms used appropriate functions (`get_state()`, `get_legal_moves()`, `make_move()`) from the game classes (see the previous section), ensuring smooth and interchangeable agent-game interactions. The unified interaction workflow simplified the experimental setup:

1. Retrieve the current state via `get_state()`.
2. Obtain available actions through `get_legal_moves()`.
3. Select the best action using algorithm-specific logic.
4. Execute the move through `make_move()`.

This clear structure facilitated direct and fair comparisons between algorithms, ensuring efficient execution of experimental scripts and reliable interpretation of results.

Experimental Context and Evaluation

Comprehensive experimental scripts (one script to compare each algorithm against the default opponent and one script to compare each algorithm against each other) were used to systematically evaluate performance under various conditions:

- **Configurations:** Flexibly configured via command-line arguments, including selection between Minimax and Q-learning, choice of alpha-beta pruning, depth limits, training episodes, and exploration parameters.
- **Training and Evaluation Cycles:** Q-learning agents were methodically trained (if chosen) over numerous episodes with periodic evaluations, clearly demonstrating learned strategic competence.
- Timeout mechanisms were used, preventing excessively long computations, especially pertinent for Minimax for connect 4 with extensive search depths (explained in the next section).

Scalability Considerations for Connect 4

Connect 4 presents severe scalability challenges for both the Minimax and Q-Learning algorithms due to the vastness of its state and action spaces.

From a theoretical standpoint, Connect 4 is played on a 6×7 grid, totaling 42 cells. A complete game can last up to 42 moves, and on each turn, a player can choose among up to 7 columns (depending on which are still available). This gives a branching factor of up to 7 and leads to a state space complexity of:

$$O(c^r) = 7^{42} \approx 1.4 \times 10^{35}$$

In comparison, Tic Tac Toe has at most $9! = 362,880$ states, a tiny fraction of Connect 4's complexity. This makes exhaustive methods like full-depth Minimax or full Q-table coverage computationally infeasible.

In practice, when I attempted to run Minimax without depth limitation for Connect 4, it entered an endless evaluation loop, even with alpha-beta pruning enabled. Early game states, in particular, have the

highest branching factor and result in an enormous game tree. To address this, I introduced the timeout configuration option as explained earlier.

The same scalability concerns affect Q-Learning. A tabular Q-Learning maintains a lookup table with entries for each (state, action) pair. While such an approach is tractable for Tic Tac Toe, it becomes unrealistic for Connect 4. Given the immense number of possible board configurations, the Q-table would require an exorbitant amount of memory and training episodes to fill meaningfully. Moreover, due to the sparse visitation of most states during training, the may fail to generalize well or converge efficiently. To make Q-Learning viable, I relied on heavy exploration during training (via ϵ -greedy action selection) and a more sophisticated heuristic for action selection during evaluation.

Evaluation

Metrics Used:

The evaluation framework employed a comprehensive set of metrics to assess both the performance and computational characteristics of the agents:

- **Game Outcome Metrics:**
 - Win/Loss/Draw rates to measure strategic effectiveness
 - Total games played to ensure statistical significance
 - Average moves per game to assess strategic efficiency
- **Computational Performance Metrics:**
 - Time per move (average, min, max) to evaluate real-time decision-making capability
 - Game duration to assess overall computational efficiency
 - States explored (for Minimax) to measure search space coverage
- **Q-Learning Specific Metrics:**
 - Episode rewards to track learning progress
 - Periodic evaluation results to measure improvement over time
 - Q-table memory usage to assess space complexity

These metrics were particularly suitable because they:

- Provide both strategic (win rates) and operational (time/space complexity) insights
- Enable direct comparisons between different types and configurations
- Capture the learning progression of Q-Learning agents
- Allow for scalability analysis through state exploration tracking
- Support real-world applicability assessment via timing measurements

The metrics were systematically collected using a dedicated `MetricsManager` class, ensuring consistent measurement and logging across all experiments. This approach enabled both detailed analysis of individual games and aggregate performance evaluation across multiple trials.

Experiment Results

The experiment scripts were run while calculating the above metrics and the following results were obtained:

Game	Match Type	Type	Player	Wins	Losses	Draws	Win Rate	Games Played	Avg. Moves	Time/Move (s)	Game Dur. (s)	States Explored
Connect 4	agent vs default	Q-Learning - Trained	Player 1	75	22	3	0.7500	100	21.2500	0.0009	0.0200	–
Connect 4	agent vs default	Minimax - with Pruning (Depth 5)	Player 1	10	0	0	1.0000	10	22.4000	0.1635	3.6600	208161
Connect 4	agent vs default	Minimax - no Pruning (Depth 5)	Player 2	10	0	0	1.0000	10	22.0000	1.2572	27.6600	1357570
Tic Tac Toe	agent vs default	Q-Learning - Untrained	Player 1	1	9	0	0.1000	10	6.5000	0.0000	0.0000	–
Tic Tac Toe	agent vs default	Q-Learning - Trained	Player 2	0	0	100	0.0000	100	9.0000	0.0000	0.0000	–
Tic Tac Toe	agent vs default	Minimax - no Pruning (Depth 5)	Player 2	0	0	10	0.0000	10	9.0000	0.0165	0.1500	86940
Tic Tac Toe	agent vs default	Minimax - with Pruning (Depth 5)	Player 2	0	0	10	0.0000	10	9.0000	0.0017	0.0200	10331
Tic Tac Toe	agent vs default	Minimax - with Pruning (Depth 5)	Player 1	0	0	10	0.0000	10	9.0000	0.0065	0.0600	36494
Tic Tac Toe	agent vs default	Minimax - with Pruning	Player 2	0	0	10	0.0000	10	9.0000	0.0014	0.0100	25484
Tic Tac Toe	agent vs default	Minimax - no Pruning and no depth	Player 2	0	0	10	0.0000	10	9.0000	0.0272	0.2500	565620

Table 1: Performance of different variants against default opponents in Tic Tac Toe and Connect 4.

Algorithms vs Default Opponents

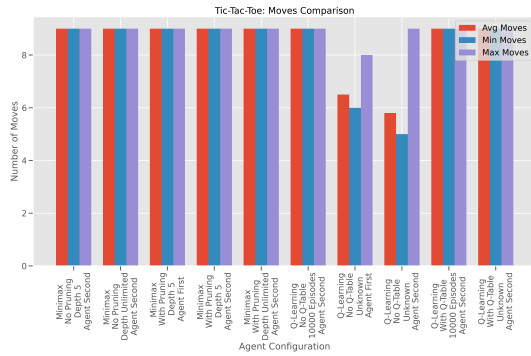


Figure 1: Moves comparison of Minimax and Q-Learning agents in Tic Tac Toe.

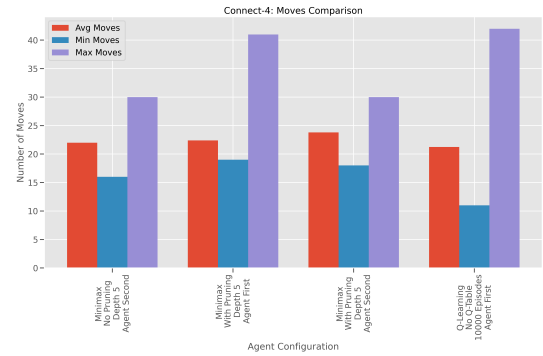


Figure 2: Moves comparison of Minimax and Q-Learning agents in Connect 4.

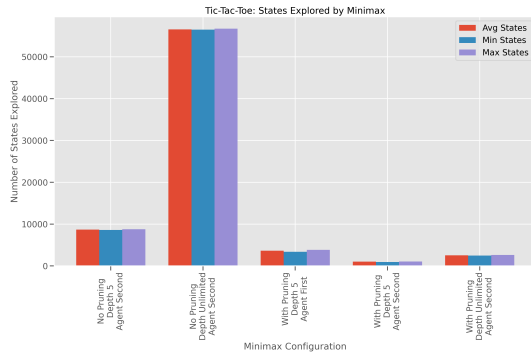


Figure 3: States explored comparison of Minimax and Q-Learning agents in Tic Tac Toe.

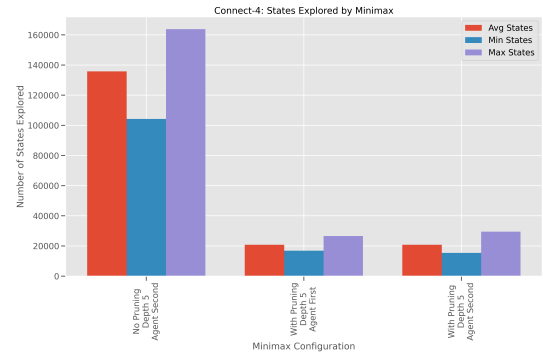


Figure 4: States explored comparison of Minimax and Q-Learning agents in Connect 4.

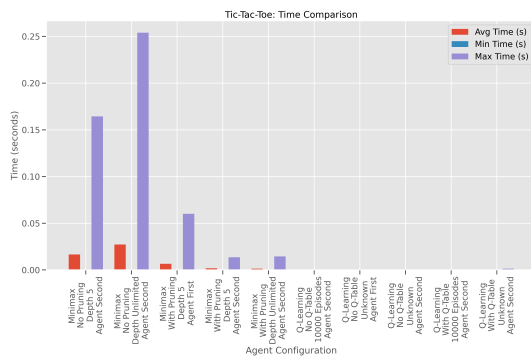


Figure 5: Time comparison of Minimax and Q-Learning agents in Tic Tac Toe.

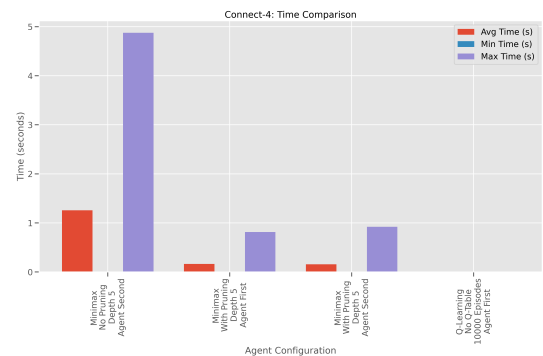


Figure 6: Time comparison of Minimax and Q-Learning agents in Connect 4.




Figure 7: Q-Learning Win Rate During Training.

Tic Tac Toe

- **Moves & States:** The results indicate that in Tic Tac Toe both algorithms perform efficiently (that is both result in draw always), generally requiring 9 moves. When we do not train a Q-table, a random is employed, leading to a variable number of moves observed in that case (see 1). Minimax typically required fewer states to explore; however, in scenarios without pruning and depth limitations, the state exploration increased significantly. This suggests that for larger games, failing to implement these strategies can lead to highly inefficient agents (see 3).
- **Decision Time:** As reflected in 5, decision times are negligible for qlearning, which is expected given that after training, its just a lookup operation. However, for minimax, the decision time is higher when there is no pruning and depth limitations, which is expected given the high branching factor of Tic Tac Toe.
- **Overall Performance:** Minimax and Q-learning agents both result in draw always, which is expected given that the default semi-intelligent opponent is very optimal for tic tac toe.

Connect 4

- **Moves & States:** For Connect 4, the average moves required increase considerably, as shown in 2. The complexity is further highlighted by the vastly higher number of states explored (refer to 4), which aligns with the theoretical state space of approximately 1.4×10^{35} configurations. When there is no pruning, the state exploration is extremely high, which is expected given the high branching factor of Connect 4.
- **Decision Time:** The decision time for Connect 4 (presented in 6) is substantially longer, especially for Minimax. For qlearning, the decision time is negligible, which is expected given that after training, its just a lookup operation. This reinforces the need for heuristics such as depth-limiting or timeout parameters to keep the search tractable.
- **Learning and Adaptation:** In the case of Q-Learning, the training win/loss rate demonstrates a gradual convergence up until 13000 episodes, after which it fluctuates around 0.75 (see 7). It suggests that the can learn upto a certain point and then it plateaus and thus requires a better method than using a q-table.
- **Overall Performance:** Minimax always wins (due to its searching capabilities), which is expected given that the default opponent is a semi-intelligent and is optimal enough for connect 4. Whereas, Q-learning is generally able to win against the default opponent but losses quite a lot of games (about 22% of the games) (see table 1) due to its final state evaluation win rate being low (around 72%).

Algorithms vs Each Other

To assess the comparative strengths of the implemented algorithms, we conducted direct head-to-head matches between Minimax and Q-Learning agents in both Tic Tac Toe and Connect 4. Each experiment consisted of 100 games, alternating player roles to ensure fairness.

The Q-Learning agents used in these evaluations were pretrained (on semi-intelligent opponent from before) but not retrained or fine-tuned against the Minimax agent. This setup helps us assess how well Q-Learning generalizes to stronger, previously unseen opponents.

Game	1 Type	2 Type	1 Wins	2 Wins	Draws	Games Played	Avg. Moves (A1/A2)	Avg. Time/Move (A1/A2, s)	States Explored	Avg. Game Duration (s)
Connect 4	Q-Learning	Minimax	0	100	0	100	8.0 / 8.0	0.0024 / 0.5008	2,212,800	4.03
Connect 4	Minimax	Q-Learning	100	0	0	100	5.0 / 4.0	0.4702 / 0.0016	1,318,500	2.36
Tic Tac Toe	Q-Learning	Minimax	0	89	11	100	3.59 / 3.48	0.0000 / 0.0055	142,830	0.02
Tic Tac Toe	Minimax	Q-Learning	43	0	57	100	4.63 / 3.63	0.0122 / 0.0000	355,791	0.06

Table 2: Comparison of Performance in Connect 4 and Tic Tac Toe

- **Minimax Consistently Outperforms Q-Learning:** Across all configurations, Minimax decisively defeats the pretrained Q-Learning agent. This reflects the ability of search-based algorithms to consistently exploit suboptimal policies learned during training.
- **Generalization Gap in Q-Learning:** The Q-Learning agents, although pretrained, were not exposed to Minimax-like adversaries during training. As a result, they fail to generalize to these stronger opponents, especially visible in the 0% win rates.
- **Efficiency and Strategy:** Q-Learning agents respond quickly with near-zero time per move, but this comes at the cost of tactical depth. Minimax agents invest more time (e.g., 0.5s per move in Connect 4) and leverage evaluation functions to identify stronger positions.
- **Game-Specific Trends:** In Tic Tac Toe, the Minimax achieves a high number of draws when playing second, reflecting optimal counterplay. In Connect 4, its dominance is absolute due to the complexity of the board and greater decision space.

Overall Comparison of Algorithms

The evaluation of Minimax and Q-Learning agents across both Tic Tac Toe and Connect 4, using both default opponents and head-to-head comparisons, reveals distinct strengths, weaknesses, and trade-offs between search-based and learning-based approaches.

- **Performance Consistency:** Minimax consistently outperformed Q-Learning in all test scenarios, especially in head-to-head matches. It was able to exploit the weaknesses of the Q-Learning agent, even when the latter was pretrained. In both games, Minimax achieved either a perfect or near-perfect win rate when facing the default opponent or the Q-Learning agent.
- **Adaptability vs Optimality:** Q-Learning showed that it can learn effective policies against specific types of opponents, particularly in Connect 4, where its pretrained model achieved a win rate of 75% against the default agent. However, it failed to generalize to stronger adversaries like Minimax, which exploited its lack of foresight and positional reasoning.
- **Resource Efficiency:** Q-Learning is significantly faster at inference time due to its table-based policy lookup. This makes it more scalable for real-time or embedded use cases. Minimax, while slower, remains tractable in practice when enhanced with alpha-beta pruning and depth limitations, even in large state spaces like Connect 4. However, the computational cost becomes significant as the search depth increases or pruning is disabled.
- **Strategic Depth:** The results confirm that Minimax with pruning and evaluation heuristics provides robust performance even under constrained resources. In contrast, Q-Learning requires extensive training and cannot adapt on-the-fly to new strategies or agents without retraining.
- **Game-Specific Observations:** - In Tic Tac Toe, both algorithms tend to converge to draws against the optimal default opponent. This reinforces the idea that the game's small state space allows even basic agents to reach near-optimal play. - In Connect 4, Minimax consistently dominated, suggesting that its tactical foresight gives it a substantial edge in deeper games. Q-Learning's performance plateaus, indicating limited policy generalization despite training.
- **Engineering Considerations:** Implementing timeouts and depth constraints was essential for ensuring that Minimax remained usable in Connect 4. Without these, it would often run indefinitely. Conversely, the Q-Learning remained responsive throughout but incurred the upfront cost of long training times and storage for the Q-table.

While Minimax is computationally more intensive, it provides reliable, optimal play when appropriately constrained. Q-Learning, on the other hand, is more scalable and efficient at runtime, but is highly dependent on training quality and lacks adaptability to unfamiliar strategies. For environments with small to moderate state spaces or where deterministic performance is essential, Minimax is the preferred choice. For large-scale, stochastic, or dynamic environments where retraining is viable, Q-Learning may be more appropriate.

Conclusion

This report highlights the journey of developing, integrating, and evaluating AI agents for two popular board games, Tic Tac Toe and Connect 4, using Minimax and Q-Learning algorithms. The designs focused on

modularity and extensibility, ensuring a clear separation of logic and interface, which facilitated smooth experimentation and integration of the agents.

The experimental results reveal that Minimax, especially when enhanced with alpha-beta pruning, delivered strong and consistent performance in both games, with particular success in Connect 4 where strategic depth plays a crucial role. On the other hand, Q-Learning proved effective in controlled scenarios and was computationally efficient during inference, but it faced challenges in generalizing against stronger opponents like Minimax, especially in unfamiliar situations.

The analysis brings to light significant trade-offs: Minimax shines in deterministic environments with manageable search spaces but can become resource-intensive in larger domains. In contrast, Q-Learning offers greater scalability during runtime and is better suited for dynamic environments, though it requires substantial training and lacks adaptability without retraining.

In summary, this project provided an engaging exploration of classical AI and reinforcement learning methods in adversarial game settings, showcasing their strengths, limitations, and practical considerations for implementation and deployment.

A Game Classes

A.1 Tic Tac Toe

```
import numpy as np
import pygame
import time
from enum import Enum

class TicTacToe:
    def __init__(self):
        self.board = np.zeros((3, 3), dtype=int) # 0 for empty, 1 for X, 2 for O
        self.current_player = 1 # Player 1 (X) starts
        self.game_over = False
        self.winner = None

    def reset(self):
        """Reset the game to initial state."""
        self.board = np.zeros((3, 3), dtype=int)
        self.current_player = 1
        self.game_over = False
        self.winner = None
        return self.get_state()

    def get_state(self):
        """Return current state of the game."""
        return self.board.copy()

    def get_legal_moves(self):
        """Return list of legal moves as (row, col) tuples."""
        if self.game_over:
            return []
        return [(i, j) for i in range(3) for j in range(3) if self.board[i, j] == 0]

    def make_move(self, move):
        """Make a move on the board.

        Args:
            move: tuple (row, col)

        Returns:
            bool: True if the move was valid, False otherwise
        """
        row, col = move
        if self.game_over or row < 0 or row > 2 or col < 0 or col > 2 or self.board[row, col] != 0:
            return False

        self.board[row, col] = self.current_player
        self._check_game_over()
        self.current_player = 3 - self.current_player # Switch players (1 -> 2, 2 -> 1)
        return True

    def _check_game_over(self):
        """Check if the game is over (win or draw)."""
        # Check rows
        for row in range(3):
            if self.board[row, 0] != 0 and self.board[row, 0] == self.board[row, 1] == self.board[row, 2]:
                self.game_over = True
                self.winner = self.board[row, 0]
                return

        # Check columns
        for col in range(3):
```

```

        if self.board[0, col] != 0 and self.board[0, col] == self.board[1, col] ==
            self.board[2, col]:
            self.game_over = True
            self.winner = self.board[0, col]
            return

    # Check diagonals
    if self.board[0, 0] != 0 and self.board[0, 0] == self.board[1, 1] == self.board[2, 2]:
        self.game_over = True
        self.winner = self.board[0, 0]
        return

    if self.board[0, 2] != 0 and self.board[0, 2] == self.board[1, 1] == self.board[2, 0]:
        self.game_over = True
        self.winner = self.board[0, 2]
        return

    # Check for draw
    if np.all(self.board != 0):
        self.game_over = True
        self.winner = 0 # Draw
        return

def is_game_over(self):
    """Return whether the game is over."""
    return self.game_over

def get_winner(self):
    """Return the winner (1 for X, 2 for O, 0 for draw, None if game not over)."""
    return self.winner

class PlayerType(Enum):
    HUMAN = 0
    AI = 1
    SEMI_INTELLIGENT = 2

class GameMode(Enum):
    HUMAN_VS_HUMAN = 0
    HUMAN_VS_AI = 1
    HUMAN_VS_SEMI = 2
    AI_VS_SEMI = 3
    AI_VS_AI = 4

class TicTacToeUI:
    def __init__(self):
        pygame.init()
        self.game = TicTacToe()
        self.width, self.height = 950, 950
        self.screen = pygame.display.set_mode((self.width, self.height))
        pygame.display.set_caption("Tic Tac Toe")

        # Colors
        self.bg_color = (240, 240, 240)
        self.line_color = (80, 80, 80)
        self.x_color = (66, 134, 244)
        self.o_color = (255, 87, 87)
        self.text_color = (50, 50, 50)
        self.highlight_color = (180, 180, 180)

        # Size and positions
        self.board_size = 450
        self.cell_size = self.board_size // 3
        self.board_margin = (self.width - self.board_size) // 2
        self.line_width = 4

        # Fonts
        self.font = pygame.font.SysFont('Arial', 30)
        self.big_font = pygame.font.SysFont('Arial', 50)

        # Game mode and players
        self.game_mode = GameMode.HUMAN_VS_HUMAN
        self.player1_type = PlayerType.HUMAN
        self.player2_type = PlayerType.HUMAN
        self.player1_agent = None # AI agent for player 1
        self.player2_agent = None # AI agent for player 2
        self.player_move = True # True if current turn is for human input
        self.ai_move_delay = 0.5 # Delay between AI moves in seconds
        self.last_ai_move_time = 0 # Last time AI made a move

    def set_game_mode(self, mode):
        """Set the game mode and initialize appropriate players."""
        self.game_mode = mode

        if mode == GameMode.HUMAN_VS_HUMAN:
            self.player1_type = PlayerType.HUMAN
            self.player2_type = PlayerType.HUMAN
            self.player_move = True

        elif mode == GameMode.HUMAN_VS_AI:
            self.player1_type = PlayerType.HUMAN
            self.player2_type = PlayerType.AI

```

```

        self.player_move = True

    elif mode == GameMode.HUMAN_VS_SEMI:
        self.player1_type = PlayerType.HUMAN
        self.player2_type = PlayerType.SEMI_INTELLIGENT
        self.player_move = True

    elif mode == GameMode.AI_VS_SEMI:
        self.player1_type = PlayerType.AI
        self.player2_type = PlayerType.SEMI_INTELLIGENT
        self.player_move = False

    elif mode == GameMode.AI_VS_AI:
        # Both players are AI agents
        self.player1_type = PlayerType.AI
        self.player2_type = PlayerType.AI
        self.player_move = False

    # Reset the game
    self.game.reset()

def set_player1_agent(self, agent):
    """Set AI agent for player 1."""
    self.player1_agent = agent

def set_player2_agent(self, agent):
    """Set AI agent for player 2."""
    self.player2_agent = agent

def run(self):
    """Main game loop."""
    running = True

    while running:
        current_time = time.time()

        # Handle events
        for event in pygame.event.get():
            if event.type == pygame.QUIT:
                running = False

            if event.type == pygame.MOUSEBUTTONDOWN:
                # Human move handling
                if self.player_move and not self.game.is_game_over():
                    self._handle_click(event.pos)

            if event.type == pygame.KEYDOWN:
                if event.key == pygame.K_r: # Reset game
                    self.game.reset()
                    self.player_move = (self.player1_type == PlayerType.HUMAN)

        # Handle AI agent moves
        if not self.game.is_game_over():
            current_player_num = self.game.current_player

            # Player 1's turn (X)
            if current_player_num == 1:
                if self.player1_type in [PlayerType.AI, PlayerType.SEMI_INTELLIGENT] and \
                    (current_time - self.last_ai_move_time >= self.ai_move_delay):
                    agent = self.player1_agent
                    if agent and hasattr(agent, 'get_move'):
                        move = agent.get_move(self.game.get_state())
                        if move:
                            self.game.make_move(move)
                            self.last_ai_move_time = current_time
                            # After AI move, it could be human's turn
                            self.player_move = (self.player2_type == PlayerType.HUMAN)

            # Player 2's turn (O)
            else:
                if self.player2_type in [PlayerType.AI, PlayerType.SEMI_INTELLIGENT] and \
                    (current_time - self.last_ai_move_time >= self.ai_move_delay):
                    agent = self.player2_agent
                    if agent and hasattr(agent, 'get_move'):
                        move = agent.get_move(self.game.get_state())
                        if move:
                            self.game.make_move(move)
                            self.last_ai_move_time = current_time
                            # After AI move, it could be human's turn
                            self.player_move = (self.player1_type == PlayerType.HUMAN)

        # Draw everything
        self._draw()
        pygame.display.flip()

        # Small delay to avoid high CPU usage
        time.sleep(0.01)

    pygame.quit()

def _handle_click(self, pos):

```

```

"""Handle mouse click to make a move."""
x, y = pos

# Check if click is within the board
if (self.board_margin <= x <= self.board_margin + self.board_size and
    self.board_margin <= y <= self.board_margin + self.board_size):

    # Convert click position to grid indices
    row = (y - self.board_margin) // self.cell_size
    col = (x - self.board_margin) // self.cell_size

    # Make move if valid
    if self.game.make_move((row, col)):
        # After human move, determine who plays next
        current_player_num = self.game.current_player
        if current_player_num == 1:
            self.player_move = (self.player1_type == PlayerType.HUMAN)
        else:
            self.player_move = (self.player2_type == PlayerType.HUMAN)

def _draw(self):
    """Draw the game board and UI."""
    # Fill background
    self.screen.fill(self.bg_color)

    # Draw title and status
    title = self.big_font.render("TicTacToe", True, self.text_color)
    self.screen.blit(title, (self.width // 2 - title.get_width() // 2, 20))

    # Display current game mode
    mode_names = {
        GameMode.HUMAN_VS_HUMAN: "Human vs Human",
        GameMode.HUMAN_VS_AI: "Human vs AI",
        GameMode.HUMAN_VS_SEMI: "Human vs Semi-Intelligent",
        GameMode.AI_VS_SEMI: "AI vs Semi-Intelligent",
        GameMode.AI_VS_AI: "AI vs AI"
    }
    mode_text = self.font.render(f"Mode: {mode_names[self.game_mode]}", True, self.text_color)
    self.screen.blit(mode_text, (self.width // 2 - mode_text.get_width() // 2, 80))

    # Game status
    if self.game.is_game_over():
        if self.game.get_winner() == 1:
            status = self.font.render("X wins!", True, self.x_color)
        elif self.game.get_winner() == 2:
            status = self.font.render("O wins!", True, self.o_color)
        else:
            status = self.font.render("Draw!", True, self.text_color)

        restart = self.font.render("Press R to restart", True, self.text_color)
        self.screen.blit(restart, (self.width // 2 - restart.get_width() // 2, self.height - 50))
    else:
        current_player_text = "X's turn" if self.game.current_player == 1 else "O's turn"
        status = self.font.render(current_player_text, True, self.x_color if
            self.game.current_player == 1 else self.o_color)

    self.screen.blit(status, (self.width // 2 - status.get_width() // 2, self.height - 100))

    # Draw board
    for i in range(4):
        # Horizontal lines
        pygame.draw.line(
            self.screen,
            self.line_color,
            (self.board_margin, self.board_margin + i * self.cell_size),
            (self.board_margin + self.board_size, self.board_margin + i * self.cell_size),
            self.line_width
        )

        # Vertical lines
        pygame.draw.line(
            self.screen,
            self.line_color,
            (self.board_margin + i * self.cell_size, self.board_margin),
            (self.board_margin + i * self.cell_size, self.board_margin + self.board_size),
            self.line_width
        )

    # Draw X's and O's
    for row in range(3):
        for col in range(3):
            center_x = self.board_margin + col * self.cell_size + self.cell_size // 2
            center_y = self.board_margin + row * self.cell_size + self.cell_size // 2

            if self.game.board[row, col] == 1: # X
                size = self.cell_size // 2 - 20
                pygame.draw.line(
                    self.screen,
                    self.x_color,
                    (center_x - size, center_y - size),

```

```

        (center_x + size, center_y + size),
        self.line_width + 3
    )
    pygame.draw.line(
        self.screen,
        self.x_color,
        (center_x - size, center_y + size),
        (center_x + size, center_y - size),
        self.line_width + 3
    )

    elif self.game.board[row, col] == 2: # 0
        size = self.cell_size // 2 - 15
        pygame.draw.circle(
            self.screen,
            self.o_color,
            (center_x, center_y),
            size,
            self.line_width + 3
        )

```

A.2 Connect 4

```

import numpy as np
import pygame
import time
from enum import Enum
from games.tic_tac_toe import PlayerType, GameMode

class Connect4:
    def __init__(self):
        self.rows = 6
        self.cols = 7
        self.board = np.zeros((self.rows, self.cols), dtype=int) # 0 for empty, 1 and 2 for
            players
        self.current_player = 1 # Player 1 starts
        self.game_over = False
        self.winner = None
        self.last_move = None

    def reset(self):
        """Reset the game to initial state."""
        self.board = np.zeros((self.rows, self.cols), dtype=int)
        self.current_player = 1
        self.game_over = False
        self.winner = None
        self.last_move = None
        return self.get_state()

    def get_state(self):
        """Return current state of the game."""
        return self.board.copy()

    def get_legal_moves(self):
        """Return list of legal moves (columns where a piece can be dropped)."""
        if self.game_over:
            return []
        return [col for col in range(self.cols) if self.board[0, col] == 0]

    def make_move(self, col):
        """Make a move by dropping a piece in the specified column.

        Args:
            col: Column to drop the piece

        Returns:
            bool: True if the move was valid, False otherwise
        """
        if self.game_over or col < 0 or col >= self.cols or self.board[0, col] != 0:
            return False

        # Find the lowest empty row in the selected column
        for row in range(self.rows - 1, -1, -1):
            if self.board[row, col] == 0:
                self.board[row, col] = self.current_player
                self.last_move = (row, col)
                break

        self._check_game_over()
        self.current_player = 3 - self.current_player # Switch players (1 -> 2, 2 -> 1)
        return True

    def _check_game_over(self):
        """Check if the game is over (win or draw)."""
        if self.last_move is None:
            return

        row, col = self.last_move

```

```

player = self.board[row, col]

# Check horizontal
for c in range(max(0, col - 3), min(col + 1, self.cols - 3)):
    if self.board[row, c] == player and self.board[row, c+1] == player \
    and self.board[row, c+2] == player and self.board[row, c+3] == player:
        self.game_over = True
        self.winner = player
        return

# Check vertical
for r in range(max(0, row - 3), min(row + 1, self.rows - 3)):
    if self.board[r, col] == player and self.board[r+1, col] == player \
    and self.board[r+2, col] == player and self.board[r+3, col] == player:
        self.game_over = True
        self.winner = player
        return

# Check diagonal (positive slope)
for r, c in zip(range(row, max(row-4, -1), -1), range(col, max(col-4, -1), -1)):
    if r+3 < self.rows and c+3 < self.cols:
        if self.board[r, c] == player and self.board[r+1, c+1] == player \
        and self.board[r+2, c+2] == player and self.board[r+3, c+3] == player:
            self.game_over = True
            self.winner = player
            return

# Check diagonal (negative slope)
for r, c in zip(range(row, max(row-4, -1), -1), range(col, min(col+4, self.cols))):
    if r+3 < self.rows and c-3 >= 0:
        if self.board[r, c] == player and self.board[r+1, c-1] == player \
        and self.board[r+2, c-2] == player and self.board[r+3, c-3] == player:
            self.game_over = True
            self.winner = player
            return

# Check for draw
if np.all(self.board[0, :] != 0):
    self.game_over = True
    self.winner = 0 # Draw
    return

def is_game_over(self):
    """Return whether the game is over."""
    return self.game_over

def get_winner(self):
    """Return the winner (1 or 2 for players, 0 for draw, None if game not over)."""
    return self.winner

class Connect4UI:
    def __init__(self):
        pygame.init()
        self.game = Connect4()
        self.cell_size = 80
        self.width = self.game.cols * self.cell_size
        self.height = (self.game.rows + 1) * self.cell_size + 200 # Extra space for UI
        self.screen = pygame.display.set_mode((self.width, self.height))
        pygame.display.set_caption("Connect4")

        # Colors
        self.bg_color = (0, 105, 148) # Blue background
        self.board_color = (0, 65, 118)
        self.player1_color = (255, 51, 51) # Red
        self.player2_color = (255, 236, 51) # Yellow
        self.text_color = (0, 0, 0)
        self.highlight_color = (0, 180, 215)

        # Fonts
        self.font = pygame.font.SysFont('Arial', 30)
        self.big_font = pygame.font.SysFont('Arial', 40)

        # Game mode and players
        self.game_mode = GameMode.HUMAN_VS_HUMAN
        self.player1_type = PlayerType.HUMAN
        self.player2_type = PlayerType.HUMAN
        self.player1_agent = None # AI agent for player 1
        self.player2_agent = None # AI agent for player 2
        self.player_move = True # True if current turn is for human input
        self.ai_move_delay = 0.5 # Delay between AI moves in seconds
        self.last_ai_move_time = 0 # Last time AI made a move

        # Mode names dictionary for UI
        self.mode_names = {
            GameMode.HUMAN_VS_HUMAN: "Human vs Human",
            GameMode.HUMAN_VS_AI: "Human vs AI",
            GameMode.HUMAN_VS_SEMI: "Human vs Semi-Intelligent",
            GameMode.AI_VS_SEMI: "AI vs Semi-Intelligent",
            GameMode.AI_VS_AI: "AI vs AI"
        }

```

```

# Animation
self.anim_active = False
self.anim_col = 0
self.anim_row = 0
self.anim_y = 0
self.anim_speed = 15
self.anim_player = 0

def set_game_mode(self, mode):
    """Set the game mode and initialize appropriate players."""
    if isinstance(mode, int):
        mode = GameMode(mode) # Convert int to GameMode Enum
    self.game_mode = mode

    if mode == GameMode.HUMAN_VS_HUMAN:
        self.player1_type = PlayerType.HUMAN
        self.player2_type = PlayerType.HUMAN
        self.player_move = True

    elif mode == GameMode.HUMAN_VS_AI:
        self.player1_type = PlayerType.HUMAN
        self.player2_type = PlayerType.AI
        self.player_move = True

    elif mode == GameMode.HUMAN_VS_SEMI:
        self.player1_type = PlayerType.HUMAN
        self.player2_type = PlayerType.SEMI_INTELLIGENT
        self.player_move = True

    elif mode == GameMode.AI_VS_SEMI:
        self.player1_type = PlayerType.AI
        self.player2_type = PlayerType.SEMI_INTELLIGENT
        self.player_move = False

    elif mode == GameMode.AI_VS_AI:
        # Both players are AI agents
        self.player1_type = PlayerType.AI
        self.player2_type = PlayerType.AI
        self.player_move = False

    # Reset the game
    self.game.reset()
    self.anim_active = False

def set_player1_agent(self, agent):
    """Set AI agent for player 1."""
    self.player1_agent = agent

def set_player2_agent(self, agent):
    """Set AI agent for player 2."""
    self.player2_agent = agent

def run(self):
    """Main game loop."""
    running = True

    while running:
        current_time = time.time()

        # Handle events
        for event in pygame.event.get():
            if event.type == pygame.QUIT:
                running = False

            if event.type == pygame.MOUSEBUTTONDOWN:
                # Human move handling
                if self.player_move and not self.game.is_game_over() and not self.anim_active:
                    self._handle_click(event.pos)

            if event.type == pygame.KEYDOWN:
                if event.key == pygame.K_r: # Reset game
                    self.game.reset()
                    self.player_move = (self.player1_type == PlayerType.HUMAN)
                    self.anim_active = False

        # Handle AI and semi-intelligent agent moves
        if not self.game.is_game_over() and not self.anim_active:
            current_player_num = self.game.current_player

            # Player 1's turn (Red)
            if current_player_num == 1:
                if self.player1_type in [PlayerType.AI, PlayerType.SEMI_INTELLIGENT] and \
                    (current_time - self.last_ai_move_time >= self.ai_move_delay):
                    agent = self.player1_agent
                    if agent and hasattr(agent, 'get_move'):
                        move = agent.get_move(self.game.get_state())
                        if move is not None:
                            self._start_animation(move)
                            self.last_ai_move_time = current_time

            # Player 2's turn (Yellow)

```



```

        else:
            if self.player2_type in [PlayerType.AI, PlayerType.SEMI_INTELLIGENT] and
               (current_time - self.last_ai_move_time >= self.ai_move_delay):
                agent = self.player2_agent
                if agent and hasattr(agent, 'get_move'):
                    move = agent.get_move(self.game.get_state())
                    if move is not None:
                        self._start_animation(move)
                        self.last_ai_move_time = current_time

# Update animation
if self.anim_active:
    self._update_animation()

# Draw everything
self._draw()
pygame.display.flip()

# Small delay to avoid high CPU usage
time.sleep(0.01)

pygame.quit()

def _start_animation(self, col):
    """Start animation for dropping a piece."""
    if col in self.game.get_legal_moves():
        self.anim_active = True
        self.anim_col = col
        self.anim_row = 0
        for row in range(self.game.rows - 1, -1, -1):
            if self.game.board[row, col] == 0:
                self.anim_row = row
                break
        self.anim_y = self.cell_size
        self.anim_player = self.game.current_player

def _update_animation(self):
    """Update dropping animation."""
    target_y = (self.anim_row + 1) * self.cell_size

    if self.anim_y < target_y:
        self.anim_y += self.anim_speed
    else:
        self.anim_active = False
        self.game.make_move(self.anim_col)

        # After a piece is dropped, determine who plays next
        current_player = self.game.current_player
        if current_player == 1:
            self.player_move = (self.player1_type == PlayerType.HUMAN)
        else:
            self.player_move = (self.player2_type == PlayerType.HUMAN)

def _handle_click(self, pos):
    """Handle mouse click to make a move."""
    x, y = pos

    # Check if click is within the valid area (above the board)
    if y < self.cell_size and 0 <= x < self.width:
        col = x // self.cell_size
        if col in self.game.get_legal_moves():
            self._start_animation(col)

def _draw(self):
    """Draw the game board and UI."""
    # Fill background
    self.screen.fill(self.bg_color)

    # Draw title and game mode info
    if self.game.is_game_over():
        if self.game.get_winner() == 1:
            status = self.big_font.render("Red wins!", True, self.player1_color)
        elif self.game.get_winner() == 2:
            status = self.big_font.render("Yellow wins!", True, self.player2_color)
        else:
            status = self.big_font.render("Draw!", True, self.text_color)

        restart = self.font.render("Press R to restart", True, self.text_color)
        self.screen.blit(restart, (self.width // 2 - restart.get_width() // 2, 20))
    else:
        if self.game.current_player == 1:
            status = self.big_font.render("Red's turn", True, self.player1_color)
        else:
            status = self.big_font.render("Yellow's turn", True, self.player2_color)

    self.screen.blit(status, (self.width // 2 - status.get_width() // 2, self.height - 160))

    # Display current game mode
    mode_text = self.font.render(f"Mode: {self.mode_names[self.game_mode]}", True,
                                self.text_color)
    self.screen.blit(mode_text, (30, self.height - 50))

```

```

# Draw board background
pygame.draw.rect(
    self.screen,
    self.board_color,
    (0, self.cell_size, self.width, self.game.rows * self.cell_size)
)

# Draw empty slots and pieces
for row in range(self.game.rows):
    for col in range(self.game.cols):
        center_x = col * self.cell_size + self.cell_size // 2
        center_y = (row + 1) * self.cell_size + self.cell_size // 2

        # Draw empty slot
        pygame.draw.circle(
            self.screen,
            self.bg_color,
            (center_x, center_y),
            self.cell_size // 2 - 5
        )

        # Draw piece
        if self.game.board[row, col] == 1:
            pygame.draw.circle(
                self.screen,
                self.player1_color,
                (center_x, center_y),
                self.cell_size // 2 - 5
            )
        elif self.game.board[row, col] == 2:
            pygame.draw.circle(
                self.screen,
                self.player2_color,
                (center_x, center_y),
                self.cell_size // 2 - 5
            )

# Draw animation
if self.anim_active:
    center_x = self.anim_col * self.cell_size + self.cell_size // 2
    center_y = self.anim_y

    color = self.player1_color if self.anim_player == 1 else self.player2_color
    pygame.draw.circle(
        self.screen,
        color,
        (center_x, center_y),
        self.cell_size // 2 - 5
    )

# Draw column highlight on hover (only when it's human's turn)
if not self.game.is_game_over() and not self.anim_active and self.player_move:
    mouse_x, mouse_y = pygame.mouse.get_pos()
    if mouse_y < self.cell_size:
        col = mouse_x // self.cell_size
        if 0 <= col < self.game.cols and col in self.game.get_legal_moves():
            pygame.draw.rect(
                self.screen,
                self.highlight_color,
                (col * self.cell_size, 0, self.cell_size, self.cell_size),
                3
            )

        # Draw preview piece
        color = self.player1_color if self.game.current_player == 1 else self.player2_color
        pygame.draw.circle(
            self.screen,
            color,
            (col * self.cell_size + self.cell_size // 2, self.cell_size // 2),
            self.cell_size // 2 - 5
        )

```

B Minimax Implementation

B.1 Minimax Base

```

import time
import numpy as np
from abc import ABC, abstractmethod

class MinimaxBase(ABC):
    """
    Abstract base class for minimax algorithm implementations.
    Game-specific implementations should inherit from this class and implement

```

```

the abstract methods.
"""

def __init__(self, max_depth=float('inf'), metrics_manager=None, use_pruning=True):
    """
    Initialize the MinimaxBase with configurable parameters.

    Args:
        max_depth (int): Maximum depth for the minimax algorithm (default: infinity)
        metrics_manager: Optional metrics manager for tracking performance
        use_pruning (bool): Whether to use alpha-beta pruning (default: True)
    """
    self.max_depth = max_depth
    self.metrics_manager = metrics_manager
    self.player = None # Will be set when get_move is called
    self.use_pruning = use_pruning
    self.states_explored = 0 # Counter for states explored

def get_move(self, state):
    """
    Get the best move for the current state using minimax algorithm.

    Args:
        state: Current state of the game

    Returns:
        The best move according to minimax algorithm
    """
    self.player = self._get_current_player(state)
    self.states_explored = 0 # Reset counter

    if self.use_pruning:
        best_move = self._find_best_move_with_pruning(state)
    else:
        best_move = self._find_best_move_without_pruning(state)

    # Record states explored if metrics manager is available
    if self.metrics_manager:
        self.metrics_manager.record_states_explored(self.states_explored)

    return best_move

def _find_best_move_with_pruning(self, state):
    """
    Find the best move using minimax algorithm with alpha-beta pruning.

    Args:
        state: Current state of the game

    Returns:
        The best move according to minimax algorithm
    """
    best_val = float('-inf')
    best_move = None
    alpha = float('-inf')
    beta = float('inf')

    for move in self._get_legal_moves(state):
        # Make the move
        new_state = self._make_move(state.copy(), move, self.player)

        # Calculate value for this move
        move_val = self._minimax_with_pruning(new_state, self.max_depth - 1, False, alpha, beta)

        # Update best move if this is better
        if move_val > best_val:
            best_val = move_val
            best_move = move

        # Update alpha
        alpha = max(alpha, best_val)

    return best_move

def _find_best_move_without_pruning(self, state):
    """
    Find the best move using minimax algorithm without alpha-beta pruning.

    Args:
        state: Current state of the game

    Returns:
        The best move according to minimax algorithm
    """
    best_val = float('-inf')
    best_move = None

    for move in self._get_legal_moves(state):
        # Make the move
        new_state = self._make_move(state.copy(), move, self.player)

```

```

        # Calculate value for this move
        move_val = self._minimax_without_pruning(new_state, self.max_depth - 1, False)

        # Update best move if this is better
        if move_val > best_val:
            best_val = move_val
            best_move = move

    return best_move

def _minimax_with_pruning(self, state, depth, is_maximizing, alpha, beta):
    """
    Minimax algorithm with alpha-beta pruning.

    Args:
        state: Current state of the game
        depth (int): Current depth in the search tree
        is_maximizing (bool): True if current player is maximizing, False otherwise
        alpha: Alpha value for pruning
        beta: Beta value for pruning

    Returns:
        The best score for the current state
    """
    # Increment state exploration counter
    self.states_explored += 1

    # Check if we've reached a terminal state
    winner = self._check_winner(state)
    if winner is not None:
        return self._evaluate_terminal(winner)

    # Check if we've reached maximum depth
    if depth == 0:
        return self._evaluate_board(state)

    # Get current player
    current_player = self._get_player(is_maximizing)

    # Maximizing player
    if is_maximizing:
        max_eval = float('-inf')
        for move in self._get_legal_moves(state):
            new_state = self._make_move(state.copy(), move, current_player)
            eval = self._minimax_with_pruning(new_state, depth - 1, False, alpha, beta)
            max_eval = max(max_eval, eval)
            alpha = max(alpha, eval)
            if beta <= alpha:
                break # Beta cutoff
        return max_eval

    # Minimizing player
    else:
        min_eval = float('inf')
        for move in self._get_legal_moves(state):
            new_state = self._make_move(state.copy(), move, current_player)
            eval = self._minimax_with_pruning(new_state, depth - 1, True, alpha, beta)
            min_eval = min(min_eval, eval)
            beta = min(beta, eval)
            if beta <= alpha:
                break # Alpha cutoff
        return min_eval

def _minimax_without_pruning(self, state, depth, is_maximizing):
    """
    Minimax algorithm without alpha-beta pruning.

    Args:
        state: Current state of the game
        depth (int): Current depth in the search tree
        is_maximizing (bool): True if current player is maximizing, False otherwise

    Returns:
        The best score for the current state
    """
    # Increment state exploration counter
    self.states_explored += 1

    # Check if we've reached a terminal state
    winner = self._check_winner(state)
    if winner is not None:
        return self._evaluate_terminal(winner)

    # Check if we've reached maximum depth
    if depth == 0:
        return self._evaluate_board(state)

    # Get current player
    current_player = self._get_player(is_maximizing)

```

```

# Get legal moves
legal_moves = self._get_legal_moves(state)

# If no legal moves, it's a draw
if not legal_moves:
    return 0

# Maximizing player
if is_maximizing:
    max_eval = float('-inf')
    for move in legal_moves:
        new_state = self._make_move(state.copy(), move, current_player)
        eval = self._minimax_without_pruning(new_state, depth - 1, False)
        max_eval = max(max_eval, eval)
    return max_eval

# Minimizing player
else:
    min_eval = float('inf')
    for move in legal_moves:
        new_state = self._make_move(state.copy(), move, current_player)
        eval = self._minimax_without_pruning(new_state, depth - 1, True)
        min_eval = min(min_eval, eval)
    return min_eval

def _get_player(self, is_maximizing):
    """
    Get the player ID based on whether it's a maximizing or minimizing move.

    Args:
        is_maximizing (bool): True if it's a maximizing move

    Returns:
        The player ID
    """
    if is_maximizing:
        return self.player
    else:
        return 3 - self.player # Switch between 1 and 2

@abstractmethod
def _get_current_player(self, state):
    """
    Get the current player from the state.

    Args:
        state: Current state of the game

    Returns:
        The current player ID
    """
    pass

@abstractmethod
def _get_legal_moves(self, state):
    """
    Get legal moves for the current state.

    Args:
        state: Current state of the game

    Returns:
        List of legal moves
    """
    pass

@abstractmethod
def _make_move(self, state, move, player):
    """
    Make a move on the board.

    Args:
        state: Current state of the game
        move: Move to make
        player: Player making the move

    Returns:
        New state after making the move
    """
    pass

@abstractmethod
def _check_winner(self, state):
    """
    Check if there's a winner in the current state.

    Args:
        state: Current state of the game

    Returns:
        The winner (1 or 2), 0 for draw, None if game is not over

```

```

    """
    pass

@abstractmethod
def _evaluate_board(self, state):
    """
    Evaluate the current board state.

    Args:
        state: Current state of the game

    Returns:
        Numerical score for the board state
    """
    pass

@abstractmethod
def _evaluate_terminal(self, winner):
    """
    Evaluate a terminal state.

    Args:
        winner: The winner (1 or 2), 0 for draw

    Returns:
        Numerical score for the terminal state
    """
    pass

```

B.2 Minimax Tic Tac Toe

```

import numpy as np
from .minimax_base import MinimaxBase

class MinimaxTicTacToe(MinimaxBase):
    """
    Minimax implementation for Tic-Tac-Toe.
    """

    def __init__(self, max_depth=9, metrics_manager=None, use_pruning=True):
        """
        Initialize the MinimaxTicTacToe with configurable parameters.
        Default max_depth is 9 since Tic-Tac-Toe has at most 9 moves.

        Args:
            max_depth (int): Maximum depth for the minimax algorithm
            metrics_manager: Optional metrics manager for tracking performance
            use_pruning (bool): Whether to use alpha-beta pruning
        """
        super().__init__(max_depth, metrics_manager, use_pruning)

    def _get_current_player(self, state):
        """
        Get the current player from the state.
        In Tic-Tac-Toe, we determine current player by counting pieces.

        Args:
            state: Current state of the game (numpy array)

        Returns:
            The current player ID (1 or 2)
        """
        # Count number of each player's pieces
        p1_count = np.count_nonzero(state == 1)
        p2_count = np.count_nonzero(state == 2)

        # Player 1 goes first, so if counts are equal, it's player 1's turn
        return 1 if p1_count <= p2_count else 2

    def _get_legal_moves(self, state):
        """
        Get legal moves for the current state.
        In Tic-Tac-Toe, legal moves are empty cells.

        Args:
            state: Current state of the game

        Returns:
            List of legal moves as (row, col) tuples
        """
        return [(i, j) for i in range(3) for j in range(3) if state[i, j] == 0]

    def _make_move(self, state, move, player):
        """
        Make a move on the board.

        Args:
            state: Current state of the game

```

```

        move: Move to make as (row, col)
        player: Player making the move

Returns:
    New state after making the move
    """
    row, col = move
    state[row, col] = player
    return state

def _check_winner(self, state):
    """
    Check if there's a winner in the current state.

    Args:
        state: Current state of the game

    Returns:
        The winner (1 or 2), 0 for draw, None if game is not over
    """
    # Check rows
    for row in range(3):
        if state[row, 0] != 0 and state[row, 0] == state[row, 1] == state[row, 2]:
            return state[row, 0]

    # Check columns
    for col in range(3):
        if state[0, col] != 0 and state[0, col] == state[1, col] == state[2, col]:
            return state[0, col]

    # Check diagonals
    if state[0, 0] != 0 and state[0, 0] == state[1, 1] == state[2, 2]:
        return state[0, 0]

    if state[0, 2] != 0 and state[0, 2] == state[1, 1] == state[2, 0]:
        return state[0, 2]

    # Check for draw (all cells filled)
    if np.all(state != 0):
        return 0 # Draw

    # Game not over yet
    return None

def _evaluate_board(self, state):
    """
    Evaluate the current board state using a heuristic.
    For Tic-Tac-Toe, we'll use a simple scoring system based on potential wins.

    Args:
        state: Current state of the game

    Returns:
        Numerical score for the board state
    """
    score = 0

    # Check rows, columns, diagonals for potential wins
    # Add points for our potential wins, subtract for opponent's

    # Check rows
    for row in range(3):
        score += self._evaluate_line(state[row, :])

    # Check columns
    for col in range(3):
        score += self._evaluate_line(state[:, col])

    # Check diagonals
    score += self._evaluate_line(np.array([state[0, 0], state[1, 1], state[2, 2]]))
    score += self._evaluate_line(np.array([state[0, 2], state[1, 1], state[2, 0]]))

    return score

def _evaluate_line(self, line):
    """
    Evaluate a line (row, column, or diagonal) for potential wins.

    Args:
        line: Array representing a line on the board

    Returns:
        Score for this line
    """
    our_pieces = np.count_nonzero(line == self.player)
    opponent_pieces = np.count_nonzero(line == (3 - self.player))
    empty_cells = np.count_nonzero(line == 0)

    # If we have a potential win (all our pieces or empty)
    if opponent_pieces == 0:
        if our_pieces == 2 and empty_cells == 1: # Almost win

```

```

        return 10
    elif our_pieces == 1 and empty_cells == 2: # Potential future win
        return 1

    # If opponent has a potential win (all their pieces or empty)
    if our_pieces == 0:
        if opponent_pieces == 2 and empty_cells == 1: # Opponent almost win
            return -10
        elif opponent_pieces == 1 and empty_cells == 2: # Opponent potential future win
            return -1

    return 0

def _evaluate_terminal(self, winner):
    """
    Evaluate a terminal state.

    Args:
        winner: The winner (1 or 2), 0 for draw

    Returns:
        Numerical score for the terminal state
    """
    if winner == 0: # Draw
        return 0
    elif winner == self.player: # We win
        return 100
    else: # Opponent wins
        return -100

```

B.3 Minimax Connect 4

```

import numpy as np
from .minimax_base import MinimaxBase

class MinimaxConnect4(MinimaxBase):
    """
    Minimax implementation for Connect-4.
    """

    def __init__(self, max_depth=4, metrics_manager=None, use_pruning=True):
        """
        Initialize the MinimaxConnect4 with configurable parameters.
        Default max_depth is 4 for Connect-4 due to its branching factor.

        Args:
            max_depth (int): Maximum depth for the minimax algorithm
            metrics_manager: Optional metrics manager for tracking performance
            use_pruning (bool): Whether to use alpha-beta pruning
        """
        super().__init__(max_depth, metrics_manager, use_pruning)
        self.rows = 6
        self.cols = 7

    def _get_current_player(self, state):
        """
        Get the current player from the state.
        In Connect-4, we determine current player by counting pieces.

        Args:
            state: Current state of the game (numpy array)

        Returns:
            The current player ID (1 or 2)
        """
        # Count number of each player's pieces
        p1_count = np.count_nonzero(state == 1)
        p2_count = np.count_nonzero(state == 2)

        # Player 1 goes first, so if counts are equal, it's player 1's turn
        return 1 if p1_count <= p2_count else 2

    def _get_legal_moves(self, state):
        """
        Get legal moves for the current state.
        In Connect-4, legal moves are columns that aren't filled.

        Args:
            state: Current state of the game

        Returns:
            List of legal moves (column indices)
        """
        return [col for col in range(self.cols) if state[0, col] == 0]

    def _make_move(self, state, move, player):
        """
        Make a move on the board.

```



```

    Args:
        state: Current state of the game
        move: Move to make (column index)
        player: Player making the move

    Returns:
        New state after making the move
    """
    col = move

    # Find the lowest empty row in the selected column
    for row in range(self.rows - 1, -1, -1):
        if state[row, col] == 0:
            state[row, col] = player
            break

    return state

def _check_winner(self, state):
    """
    Check if there's a winner in the current state.

    Args:
        state: Current state of the game

    Returns:
        The winner (1 or 2), 0 for draw, None if game is not over
    """
    # Check horizontal
    for row in range(self.rows):
        for col in range(self.cols - 3):
            if state[row, col] != 0 and state[row, col] == state[row, col+1] == state[row, col+2] == state[row, col+3]:
                return state[row, col]

    # Check vertical
    for row in range(self.rows - 3):
        for col in range(self.cols):
            if state[row, col] != 0 and state[row, col] == state[row+1, col] == state[row+2, col] == state[row+3, col]:
                return state[row, col]

    # Check diagonal (positive slope)
    for row in range(self.rows - 3):
        for col in range(self.cols - 3):
            if state[row, col] != 0 and state[row, col] == state[row+1, col+1] == state[row+2, col+2] == state[row+3, col+3]:
                return state[row, col]

    # Check diagonal (negative slope)
    for row in range(3, self.rows):
        for col in range(self.cols - 3):
            if state[row, col] != 0 and state[row, col] == state[row-1, col+1] == state[row-2, col+2] == state[row-3, col+3]:
                return state[row, col]

    # Check for draw (top row filled)
    if np.all(state[0, :] != 0):
        return 0 # Draw

    # Game not over yet
    return None

def _count_immediate_threats(self, state):
    """
    Count the number of immediate winning threats for the current player.
    An immediate threat is an empty position that would result in a win.

    Args:
        state: Current state of the game

    Returns:
        Number of immediate threats
    """
    threats = 0

    # Try each possible move
    for col in range(self.cols):
        if state[0, col] != 0: # Column is full
            continue

        # Find where piece would land
        for row in range(self.rows-1, -1, -1):
            if state[row, col] == 0:
                # Try the move
                test_state = state.copy()
                test_state[row, col] = self.player

                # Check if this creates a win
                if self._is_winning_move(test_state, row, col):
                    threats += 1

```

```

        break

    return threats

def _is_winning_move(self, state, row, col):
    """
    Check if the last move at (row, col) creates a win.
    More efficient than checking the entire board.

    Args:
        state: Current state of the game
        row: Row of last move
        col: Column of last move

    Returns:
        True if the move creates a win
    """
    player = state[row, col]

    # Check horizontal
    count = 0
    for c in range(max(0, col-3), min(self.cols, col+4)):
        if state[row, c] == player:
            count += 1
            if count == 4:
                return True
        else:
            count = 0

    # Check vertical
    count = 0
    for r in range(max(0, row-3), min(self.rows, row+4)):
        if state[r, col] == player:
            count += 1
            if count == 4:
                return True
        else:
            count = 0

    # Check diagonal (positive slope)
    count = 0
    for i in range(-3, 4):
        r = row + i
        c = col + i
        if 0 <= r < self.rows and 0 <= c < self.cols:
            if state[r, c] == player:
                count += 1
                if count == 4:
                    return True
            else:
                count = 0

    # Check diagonal (negative slope)
    count = 0
    for i in range(-3, 4):
        r = row - i
        c = col + i
        if 0 <= r < self.rows and 0 <= c < self.cols:
            if state[r, c] == player:
                count += 1
                if count == 4:
                    return True
            else:
                count = 0

    return False

def _evaluate_board(self, state):
    """
    Evaluate the current board state using a heuristic.
    For Connect-4, we'll use a weighted scoring system based on potential connections.

    Args:
        state: Current state of the game

    Returns:
        Numerical score for the board state
    """
    score = 0

    # Check all possible four-in-a-row windows and score them

    # Horizontal windows
    for row in range(self.rows):
        for col in range(self.cols - 3):
            window = state[row, col:col+4]
            score += self._evaluate_window(window)

    # Vertical windows
    for row in range(self.rows - 3):
        for col in range(self.cols):

```

```

        window = state[row:row+4, col]
        score += self._evaluate_window(window)

# Positive diagonal windows
for row in range(self.rows - 3):
    for col in range(self.cols - 3):
        window = np.array([state[row+i, col+i] for i in range(4)])
        score += self._evaluate_window(window)

# Negative diagonal windows
for row in range(3, self.rows):
    for col in range(self.cols - 3):
        window = np.array([state[row-i, col+i] for i in range(4)])
        score += self._evaluate_window(window)

# Prefer center columns (better positions strategically)
center_col = self.cols // 2
for row in range(self.rows):
    if state[row, center_col] == self.player:
        # More weight to lower positions
        score += 3 * (self.rows - row)

# Vertical threat bonus (they're harder to block)
for col in range(self.cols):
    for row in range(self.rows - 3):
        window = state[row:row+4, col]
        our_pieces = np.count_nonzero(window == self.player)
        empty_slots = np.count_nonzero(window == 0)
        if our_pieces == 3 and empty_slots == 1:
            score += 2 # Additional bonus for vertical threats

# Detect forced moves (immediate threats)
immediate_threats = self._count_immediate_threats(state)
if immediate_threats > 1:
    score += 50 # Multiple threats usually lead to forced win

return score

def _evaluate_window(self, window):
    """
    Evaluate a window of 4 slots for potential connections.

    Args:
        window: Array of 4 positions

    Returns:
        Score for this window
    """
    our_pieces = np.count_nonzero(window == self.player)
    opponent_pieces = np.count_nonzero(window == (3 - self.player))
    empty_slots = np.count_nonzero(window == 0)

    if our_pieces == 4:
        return 100 # We win
    elif our_pieces == 3 and empty_slots == 1:
        return 5 # Potential win next move
    elif our_pieces == 2 and empty_slots == 2:
        return 2 # Potential future win

    if opponent_pieces == 3 and empty_slots == 1:
        return -4 # Block opponent win

    return 0

def _evaluate_terminal(self, winner):
    """
    Evaluate a terminal state.

    Args:
        winner: The winner (1 or 2), 0 for draw

    Returns:
        Numerical score for the terminal state
    """
    if winner == 0: # Draw
        return 0
    elif winner == self.player: # We win
        return 1000000 # Very high value
    else: # Opponent wins
        return -1000000 # Very low value

```

C Q-Learning Implementation

C.1 Q-Learning Base

```
import numpy as np
```

```

import random
import pickle
import os
from abc import ABC, abstractmethod

class QLearningAgent(ABC):
    def __init__(self, player_number=1, alpha=0.1, gamma=0.9, epsilon=0.1,
                  epsilon_decay=0.999, epsilon_min=0.01, metrics_manager=None):
        """
        Initialize the Q-learning agent.

        Args:
            player_number: The player number (1 or 2)
            alpha: Learning rate
            gamma: Discount factor
            epsilon: Exploration rate
            epsilon_decay: Rate at which epsilon decays
            epsilon_min: Minimum exploration rate
            metrics_manager: Metrics manager for tracking stats
        """
        self.player_number = player_number
        self.alpha = alpha
        self.gamma = gamma
        self.epsilon = epsilon
        self.epsilon_decay = epsilon_decay
        self.epsilon_min = epsilon_min
        self.metrics_manager = metrics_manager

        # Initialize Q-table
        self.q_table = {}

        # Set rewards
        self.reward_win = 1.0
        self.reward_loss = -1.0
        self.reward_draw = 0.0
        self.reward_move = -0.01 # Small negative reward for each move to encourage faster winning

        # Training history
        self.training_stats = {
            'episode_rewards': [],
            'win_rate': [],
            'episode_lengths': []
        }

    def decay_epsilon(self):
        """Decay the exploration rate."""
        if self.epsilon > self.epsilon_min:
            self.epsilon *= self.epsilon_decay

    def get_q_value(self, state, action):
        """Get Q-value for state-action pair. Return 0 if not visited before."""
        state_key = self.state_to_key(state)
        if state_key in self.q_table and action in self.q_table[state_key]:
            return self.q_table[state_key][action]
        return 0.0

    def update_q_value(self, state, action, reward, next_state):
        """Update Q-value for state-action pair."""
        state_key = self.state_to_key(state)
        next_state_key = self.state_to_key(next_state)

        # Initialize q_table entry if it doesn't exist
        if state_key not in self.q_table:
            self.q_table[state_key] = {}

        # Get current Q-value
        current_q = self.get_q_value(state, action)

        # Get max Q-value for next state
        next_max_q = 0.0
        if next_state_key in self.q_table:
            next_q_values = self.q_table[next_state_key].values()
            if next_q_values:
                next_max_q = max(next_q_values)

        # Update rule: Q(s,a) = Q(s,a) + alpha * (r + gamma * max(Q(s',a')) - Q(s,a))
        new_q = current_q + self.alpha * (reward + self.gamma * next_max_q - current_q)
        self.q_table[state_key][action] = new_q

    def choose_action(self, state, legal_moves, training=False):
        """Choose an action using epsilon-greedy policy."""
        if training and random.random() < self.epsilon:
            # Exploration: choose a random action
            return random.choice(legal_moves) if legal_moves else None

        # Exploitation: choose the best action
        best_action = None
        best_value = float('-inf')

        # Shuffle the legal moves to break ties randomly
        random.shuffle(legal_moves)

```

```

for action in legal_moves:
    action_value = self.get_q_value(state, action)
    if action_value > best_value:
        best_value = action_value
        best_action = action

return best_action

def get_move(self, state):
    """Get a move for the current state (used during gameplay)."""
    legal_moves = self.get_legal_moves(state)
    if not legal_moves:
        return None
    return self.choose_action(state, legal_moves, training=False)

def train(self, num_episodes=10000, eval_interval=500, eval_games=50, opponent=None):
    """Train the agent through self-play or against an opponent."""
    total_rewards = []

    # Check if metrics manager is available
    if self.metrics_manager:
        self.metrics_manager.set_q_table(self.q_table)

    for episode in range(1, num_episodes + 1):
        # Reset the game
        game = self.create_game()
        state = game.get_state()
        done = False
        episode_reward = 0
        moves = 0

        while not done:
            # Get legal moves
            legal_moves = game.get_legal_moves()
            if not legal_moves:
                break

            # Choose action
            action = self.choose_action(state, legal_moves, training=True)

            # Make the move
            game.make_move(action)
            next_state = game.get_state()
            moves += 1

            # Check if game is over
            if game.is_game_over():
                winner = game.get_winner()
                if winner == self.player_number:
                    reward = self.reward_win
                elif winner == 0: # Draw
                    reward = self.reward_draw
                else: # Loss
                    reward = self.reward_loss
                done = True
            else:
                reward = self.reward_move

            # If playing against an opponent, let opponent make a move
            if opponent and not done:
                opponent_move = opponent.get_move(game.get_state())
                if opponent_move:
                    game.make_move(opponent_move)
                    # Check if opponent won
                    if game.is_game_over():
                        winner = game.get_winner()
                        if winner == 3 - self.player_number: # Opponent won
                            reward = self.reward_loss
                        elif winner == 0: # Draw
                            reward = self.reward_draw
                        done = True
                    next_state = game.get_state()
                else:
                    done = True

            # Update Q-values
            self.update_q_value(state, action, reward, next_state)

            # Update state
            state = next_state
            episode_reward += reward

        # Decay exploration rate
        self.decay_epsilon()

        # Record stats
        total_rewards.append(episode_reward)
        self.training_stats['episode_rewards'].append(episode_reward)
        self.training_stats['episode_lengths'].append(moves)

```

```

        # Log with metrics manager
        if self.metrics_manager:
            self.metrics_manager.record_q_learning_reward(episode, episode_reward)

    # Periodically evaluate the agent
    if episode % eval_interval == 0:
        win_rate = self.evaluate(eval_games, opponent)
        self.training_stats['win_rate'].append((episode, win_rate))

        if self.metrics_manager:
            self.metrics_manager.print_q_table_memory()

        print(f"Episode_{episode}/{num_episodes}: Win rate: {win_rate:.2f}, "
              f"Epsilon: {self.epsilon:.3f}, "
              f"Q-table size: {len(self.q_table)}")

    return self.training_stats

def evaluate(self, num_games=100, opponent=None):
    """Evaluate the agent by playing against an opponent or randomly."""
    results = []

    for _ in range(num_games):
        game = self.create_game()
        done = False

        while not done:
            # Agent's turn
            if game.current_player == self.player_number:
                action = self.get_move(game.get_state())
                if action is None:
                    break
                game.make_move(action)
            # Opponent's turn
            else:
                if opponent:
                    opp_action = opponent.get_move(game.get_state())
                else:
                    # Random opponent
                    legal_moves = game.get_legal_moves()
                    opp_action = random.choice(legal_moves) if legal_moves else None

                if opp_action is None:
                    break
                game.make_move(opp_action)

            # Check if game is over
            if game.is_game_over():
                winner = game.get_winner()
                if winner == self.player_number:
                    results.append('win')
                elif winner == 0:
                    results.append('draw')
                else:
                    results.append('loss')
                done = True

    # Record evaluation results with metrics manager
    if self.metrics_manager:
        self.metrics_manager.record_q_learning_evaluation(len(self.training_stats['episode_rewards']),
                                                         results)

    win_rate = results.count('win') / len(results) if results else 0
    return win_rate

def save(self, filepath):
    """Save the Q-table to a file."""
    with open(filepath, 'wb') as f:
        pickle.dump(self.q_table, f)

    # Also save training stats
    stats_filepath = f"{os.path.splitext(filepath)[0]}_stats.pkl"
    with open(stats_filepath, 'wb') as f:
        pickle.dump(self.training_stats, f)

def load(self, filepath):
    """Load the Q-table from a file."""
    with open(filepath, 'rb') as f:
        self.q_table = pickle.load(f)

    # Also try to load training stats
    stats_filepath = f"{os.path.splitext(filepath)[0]}_stats.pkl"
    if os.path.exists(stats_filepath):
        with open(stats_filepath, 'rb') as f:
            self.training_stats = pickle.load(f)

    # Update metrics manager
    if self.metrics_manager:
        self.metrics_manager.set_q_table(self.q_table)

```

@abstractmethod

```

def state_to_key(self, state):
    """Convert state to a key that can be used in the Q-table.
    Must be implemented by subclasses."""
    pass

@abstractmethod
def get_legal_moves(self, state):
    """Get legal moves for the given state.
    Must be implemented by subclasses."""
    pass

@abstractmethod
def create_game(self):
    """Create a new game instance.
    Must be implemented by subclasses."""
    pass

```

C.2 Q-Learning Tic Tac Toe

```

import numpy as np
from games.tic_tac_toe import TicTacToe
from .qlearning_base import QLearningAgent

class QLearningTicTacToe(QLearningAgent):
    def __init__(self, player_number=1, alpha=0.3, gamma=0.9, epsilon=0.3,
                  epsilon_decay=0.9999, epsilon_min=0.01, metrics_manager=None):
        """
        Initialize the Q-learning agent for Tic-Tac-Toe.

        Args:
            player_number: The player number (1 or 2)
            alpha: Learning rate
            gamma: Discount factor
            epsilon: Exploration rate
            epsilon_decay: Rate at which epsilon decays
            epsilon_min: Minimum exploration rate
            metrics_manager: Metrics manager for tracking stats
        """
        super().__init__(player_number, alpha, gamma, epsilon,
                          epsilon_decay, epsilon_min, metrics_manager)

        # Higher rewards for Tic-Tac-Toe due to shorter game length
        self.reward_win = 1.0
        self.reward_loss = -1.0
        self.reward_draw = 0.2 # Draws are better than losses
        self.reward_move = -0.05 # Small penalty for each move

    def state_to_key(self, state):
        """
        Convert the game state to a hashable key for the Q-table.

        Args:
            state: 3x3 numpy array representing the game board

        Returns:
            tuple: A hashable representation of the state
        """
        # Convert the board state to a tuple of tuples
        # This ensures the state is hashable and can be used as a dictionary key
        return tuple(map(tuple, state))

    def get_legal_moves(self, state):
        """
        Get all legal moves for the current state.

        Args:
            state: The game state

        Returns:
            list: List of legal moves as (row, col) tuples
        """
        return [(i, j) for i in range(3) for j in range(3) if state[i, j] == 0]

    def create_game(self):
        """Create a new Tic-Tac-Toe game instance."""
        game = TicTacToe()
        # If player 2, make a random first move as player 1
        if self.player_number == 2:
            # Choose a random move for player 1
            moves = self.get_legal_moves(game.get_state())
            if moves:
                move = moves[np.random.randint(len(moves))]
                game.make_move(move)
        return game

    def get_symmetries(self, state, action):
        """
        Get all symmetric states and corresponding actions.

```

This helps the agent learn faster by exploiting the symmetry of the game.

```

Args:
    state: The game state
    action: The action taken in that state

Returns:
    list: List of (state_key, action) pairs for all symmetries
"""
state_array = np.array(state).reshape(3, 3)
row, col = action

symmetries = []

# Original
symmetries.append((tuple(map(tuple, state_array)), (row, col)))

# Rotate 90 degrees
rot90 = np.rot90(state_array)
new_row, new_col = 2 - col, row
symmetries.append((tuple(map(tuple, rot90)), (new_row, new_col)))

# Rotate 180 degrees
rot180 = np.rot90(rot90)
new_row, new_col = 2 - row, 2 - col
symmetries.append((tuple(map(tuple, rot180)), (new_row, new_col)))

# Rotate 270 degrees
rot270 = np.rot90(rot180)
new_row, new_col = col, 2 - row
symmetries.append((tuple(map(tuple, rot270)), (new_row, new_col)))

# Flip horizontally
flip_h = np.fliplr(state_array)
new_row, new_col = row, 2 - col
symmetries.append((tuple(map(tuple, flip_h)), (new_row, new_col)))

# Flip vertically
flip_v = np.flipud(state_array)
new_row, new_col = 2 - row, col
symmetries.append((tuple(map(tuple, flip_v)), (new_row, new_col)))

# Flip along main diagonal
flip_diag = np.transpose(state_array)
new_row, new_col = col, row
symmetries.append((tuple(map(tuple, flip_diag)), (new_row, new_col)))

# Flip along other diagonal
flip_diag2 = np.rot90(np.transpose(rot90))
new_row, new_col = 2 - col, 2 - row
symmetries.append((tuple(map(tuple, flip_diag2)), (new_row, new_col)))

return symmetries

def update_q_value(self, state, action, reward, next_state):
    """
    Update Q-value for state-action pair and all its symmetries.

    Args:
        state: The game state
        action: The action taken
        reward: The reward received
        next_state: The resulting state
    """
    # Get all symmetric states and actions
    symmetries = self.get_symmetries(state, action)

    for sym_state, sym_action in symmetries:
        state_key = sym_state

        # Initialize q_table entry if it doesn't exist
        if state_key not in self.q_table:
            self.q_table[state_key] = {}

        # Get current Q-value
        if sym_action in self.q_table[state_key]:
            current_q = self.q_table[state_key][sym_action]
        else:
            current_q = 0.0

        # Get max Q-value for next state
        next_state_key = self.state_to_key(next_state)
        next_max_q = 0.0
        if next_state_key in self.q_table:
            next_q_values = self.q_table[next_state_key].values()
            if next_q_values:
                next_max_q = max(next_q_values)

        # Update rule: Q(s,a) = Q(s,a) + alpha * (r + gamma * max(Q(s',a')) - Q(s,a))
        new_q = current_q + self.alpha * (reward + self.gamma * next_max_q - current_q)
        self.q_table[state_key][sym_action] = new_q

```


C.3 Q-Learning Connect 4

```
import random
import numpy as np
from games.connect4 import Connect4
from .qlearning_base import QLearningAgent

class QLearningConnect4(QLearningAgent):
    def __init__(self, player_number=1, alpha=0.2, gamma=0.95, epsilon=0.3,
                  epsilon_decay=0.9995, epsilon_min=0.01, metrics_manager=None):
        """
        Initialize the Q-learning agent for Connect-4.

        Args:
            player_number: The player number (1 or 2)
            alpha: Learning rate
            gamma: Discount factor
            epsilon: Exploration rate
            epsilon_decay: Rate at which epsilon decays
            epsilon_min: Minimum exploration rate
            metrics_manager: Metrics manager for tracking stats
        """
        super().__init__(player_number, alpha, gamma, epsilon,
                          epsilon_decay, epsilon_min, metrics_manager)

        # Adjust rewards for Connect-4
        self.reward_win = 1.0
        self.reward_loss = -1.0
        self.reward_draw = 0.0
        self.reward_move = -0.01 # Small penalty for each move

        # Caching for move detection
        self._horizontal_window_indices = self._precompute_horizontal_windows()
        self._vertical_window_indices = self._precompute_vertical_windows()
        self._diagonal_window_indices = self._precompute_diagonal_windows()

    def state_to_key(self, state):
        """
        Convert the game state to a hashable key for the Q-table.
        For Connect-4, we use a tuple of tuples representation.

        Args:
            state: 6x7 numpy array representing the game board

        Returns:
            tuple: A hashable representation of the state
        """
        return tuple(map(tuple, state))

    def get_legal_moves(self, state):
        """
        Get all legal moves for the current state in Connect-4.
        Legal moves are columns where a piece can be dropped.

        Args:
            state: The game state

        Returns:
            list: List of legal moves (column indices)
        """
        return [col for col in range(7) if state[0][col] == 0]

    def create_game(self):
        """Create a new Connect-4 game instance."""
        game = Connect4()
        # If player 2, make a random first move as player 1
        if self.player_number == 2:
            # Choose a random move for player 1
            moves = self.get_legal_moves(game.get_state())
            if moves:
                move = moves[np.random.randint(len(moves))]
                game.make_move(move)
        return game

    def _precompute_horizontal_windows(self):
        """Precompute all possible horizontal 4-in-a-row window indices."""
        windows = []
        for row in range(6):
            for col in range(4):
                window = [(row, col + i) for i in range(4)]
                windows.append(window)
        return windows

    def _precompute_vertical_windows(self):
        """Precompute all possible vertical 4-in-a-row window indices."""
        windows = []
```

```

    for row in range(3):
        for col in range(7):
            window = [(row + i, col) for i in range(4)]
            windows.append(window)
    return windows

def _precompute_diagonal_windows(self):
    """Precompute all possible diagonal 4-in-a-row window indices."""
    windows = []
    # Positive slope diagonals
    for row in range(3, 6):
        for col in range(4):
            window = [(row - i, col + i) for i in range(4)]
            windows.append(window)
    # Negative slope diagonals
    for row in range(3):
        for col in range(4):
            window = [(row + i, col + i) for i in range(4)]
            windows.append(window)
    return windows

def _count_window(self, state, window, player):
    """
    Count pieces in a window for a given player.
    Returns the count of player's pieces if the window doesn't contain opponent pieces,
    otherwise returns 0.
    """
    count = 0
    opponent = 3 - player
    for row, col in window:
        if state[row][col] == opponent:
            return 0
        if state[row][col] == player:
            count += 1
    return count

def _detect_threats(self, state, player):
    """
    Detect immediate threats (3-in-a-row with an empty space) for the given player.

    Args:
        state: The game state
        player: The player to check threats for

    Returns:
        list: List of column indices where there are immediate threats
    """
    threats = []
    opponent = 3 - player

    # Check horizontal threats
    for window in self._horizontal_window_indices:
        # Count player and empty spaces in window
        player_count = 0
        empty_pos = None

        for row, col in window:
            if state[row][col] == player:
                player_count += 1
            elif state[row][col] == 0:
                empty_pos = (row, col)

        # If 3 player pieces and 1 empty, it's a threat
        if player_count == 3 and empty_pos:
            # Make sure the empty position is valid (either at bottom or has support below)
            empty_row, empty_col = empty_pos
            if empty_row == 5 or state[empty_row + 1][empty_col] != 0:
                if empty_col not in threats:
                    threats.append(empty_col)

    # Check vertical threats
    for window in self._vertical_window_indices:
        # For vertical windows, we need 3 player pieces and the top is empty
        cells = [(r, c) for r, c in window]
        player_count = sum(1 for r, c in cells if state[r][c] == player)
        bottom_cell = max(cells, key=lambda x: x[0]) # Cell with largest row index
        top_cell = min(cells, key=lambda x: x[0]) # Cell with smallest row index

        if player_count == 3 and state[top_cell[0]][top_cell[1]] == 0:
            if top_cell[1] not in threats:
                threats.append(top_cell[1])

    # Check diagonal threats (both directions)
    for window in self._diagonal_window_indices:
        player_count = 0
        empty_pos = None

        for row, col in window:
            if state[row][col] == player:
                player_count += 1
            elif state[row][col] == 0:

```

```

        empty_pos = (row, col)

        # If 3 player pieces and 1 empty, check if it's a valid move
        if player_count == 3 and empty_pos:
            empty_row, empty_col = empty_pos
            # Check if the move is valid (either at bottom or has support)
            if empty_row == 5 or state[empty_row + 1][empty_col] != 0:
                if empty_col not in threats:
                    threats.append(empty_col)

    return threats

def _detect_double_threats(self, state, player):
    """
    Detect positions that would create multiple threats simultaneously.
    These are usually winning moves.

    Args:
        state: The game state
        player: The player to check threats for

    Returns:
        list: List of column indices that create multiple threats
    """
    double_threats = []

    # For each legal move, simulate it and count resulting threats
    for col in range(7):
        # Skip if column is full
        if state[0][col] != 0:
            continue

        # Find the row where the piece would land
        for row in range(5, -1, -1):
            if state[row][col] == 0:
                # Simulate the move
                temp_state = [list(row) for row in state]
                temp_state[row][col] = player

                # Count threats after this move
                threats = self._detect_threats(temp_state, player)
                if len(threats) >= 2:
                    double_threats.append(col)
                    break

    return double_threats

def get_heuristic_features(self, state):
    """
    Extract heuristic features from the state to augment the Q-learning.
    Enhanced with strategies from minimax evaluation function.

    Features returned:
    - Number of potential winning lines with 1, 2, or 3 pieces
    - Center column control with positional weighting
    - Vertical threat recognition
    - Multiple threats detection

    Args:
        state: The game state

    Returns:
        dict: Dictionary of features
    """
    features = {
        'one_piece': 0,
        'two_pieces': 0,
        'three_pieces': 0,
        'center_control': 0,
        'vertical_threats': 0,
        'has_immediate_threat': 0,
        'has_double_threat': 0,
        'blocking_opponent_win': 0
    }

    player = self.player_number
    opponent = 3 - player

    # Count pieces in horizontal windows
    for window in self._horizontal_window_indices:
        count = self._count_window(state, window, player)
        if count == 1:
            features['one_piece'] += 1
        elif count == 2:
            features['two_pieces'] += 1
        elif count == 3:
            features['three_pieces'] += 1
        # Check if this is a valid threat (can be played immediately)
        for row, col in window:
            if state[row][col] == 0:
                # If this empty cell is at bottom or has support

```

```

        if row == 5 or state[row+1][col] != 0:
            features['has_immediate_threat'] = 1

# Count pieces in vertical windows
for window in self._vertical_window_indices:
    count = self._count_window(state, window, player)
    if count == 1:
        features['one_piece'] += 1
    elif count == 2:
        features['two_pieces'] += 1
    elif count == 3:
        features['three_pieces'] += 1
        features['vertical_threats'] += 1
        features['has_immediate_threat'] = 1

# Count pieces in diagonal windows
for window in self._diagonal_window_indices:
    count = self._count_window(state, window, player)
    if count == 1:
        features['one_piece'] += 1
    elif count == 2:
        features['two_pieces'] += 1
    elif count == 3:
        features['three_pieces'] += 1
        # Check if this is a valid threat (can be played immediately)
        for row, col in window:
            if state[row][col] == 0:
                # If this empty cell is at bottom or has support
                if row == 5 or state[row+1][col] != 0:
                    features['has_immediate_threat'] = 1

# Center column control with positional weighting
center_col = 3
for row in range(6):
    if state[row][center_col] == player:
        # More weight to lower positions (like in minimax)
        features['center_control'] += (6 - row)

# Check for multiple threats
threats = self._detect_threats(state, player)
if len(threats) >= 2:
    features['has_double_threat'] = 1

# Check if we need to block opponent's immediate win
opponent_threats = self._detect_threats(state, opponent)
if opponent_threats:
    features['blocking_opponent_win'] = 1

return features

def choose_action(self, state, legal_moves, training=False):
    """
    Choose an action using epsilon-greedy policy combined with enhanced heuristic knowledge.
    This implementation gives a sophisticated evaluation of potential moves.

    Args:
        state: The game state
        legal_moves: List of legal moves
        training: Whether we're in training mode (exploration enabled)

    Returns:
        int: Column to drop the piece
    """
    if not legal_moves:
        return None

    # During training, use epsilon-greedy exploration
    if training and random.random() < self.epsilon:
        return random.choice(legal_moves)

    # Check if we have an immediate winning move
    player_threats = self._detect_threats(state, self.player_number)
    if player_threats:
        # Prioritize the winning move
        return player_threats[0]

    # Check if we need to block opponent's immediate win
    opponent = 3 - self.player_number
    opponent_threats = self._detect_threats(state, opponent)
    if opponent_threats:
        # Need to block
        return opponent_threats[0]

    # Check for moves that create multiple threats (usually winning)
    double_threats = self._detect_double_threats(state, self.player_number)
    if double_threats:
        return double_threats[0]

    # Exploitation: choose the best action based on Q-values and enhanced heuristics
    best_action = None
    best_value = float('-inf')

```

```

# Weight of heuristic values vs Q-values
heuristic_weight = 0.3

for action in legal_moves:
    # Get Q-value
    q_value = self.get_q_value(state, action)

    # Simulate the move to get heuristic value
    temp_state = np.array(state)
    for row in range(5, -1, -1):
        if temp_state[row][action] == 0:
            temp_state[row][action] = self.player_number
            break

    temp_features = self.get_heuristic_features(temp_state)

    # Calculate heuristic value with improved weights
    heuristic_value = (
        0.1 * temp_features['one_piece'] +
        0.3 * temp_features['two_pieces'] +
        0.8 * temp_features['three_pieces'] +
        0.4 * temp_features['center_control'] +
        0.5 * temp_features['vertical_threats'] +
        2.0 * temp_features['has_immediate_threat'] +
        3.0 * temp_features['has_double_threat'] +
        1.0 * temp_features['blocking_opponent_win']
    )

    # Combine Q-value and heuristic
    total_value = (1 - heuristic_weight) * q_value + heuristic_weight * heuristic_value

    if total_value > best_value:
        best_value = total_value
        best_action = action

return best_action

def update_q_value(self, state, action, reward, next_state):
    """
    Update Q-value for state-action pair.
    Connect-4 has too many states for symmetry-based optimization to be practical.

    Args:
        state: The game state
        action: The action taken
        reward: The reward received
        next_state: The resulting state
    """
    state_key = self.state_to_key(state)
    next_state_key = self.state_to_key(next_state)

    # Initialize q_table entry if it doesn't exist
    if state_key not in self.q_table:
        self.q_table[state_key] = {}

    # Get current Q-value
    if action in self.q_table[state_key]:
        current_q = self.q_table[state_key][action]
    else:
        current_q = 0.0

    # Get max Q-value for next state
    next_max_q = 0.0
    if next_state_key in self.q_table:
        next_q_values = self.q_table[next_state_key].values()
        if next_q_values:
            next_max_q = max(next_q_values)

    # Update rule:  $Q(s,a) = Q(s,a) + \alpha * (r + \gamma * \max_{a'} Q(s',a') - Q(s,a))$ 
    new_q = current_q + self.alpha * (reward + self.gamma * next_max_q - current_q)
    self.q_table[state_key][action] = new_q

```

D Default Opponent Implementation

D.1 Default Opponent Tic Tac Toe

```

import random
import numpy as np

from games.tic_tac_toe import TicTacToe

class DefaultOpponentTTT:
    """A semi-intelligent agent for Tic Tac Toe that:
    - Plays a winning move if available
    - Blocks opponent's winning move if possible
    """

```

```

- Otherwise plays randomly
"""
def __init__(self, player_number=2):
    """Initialize the agent.

    Args:
        player_number: 1 for X, 2 for O (default is 2)
    """
    self.player_number = player_number

def get_move(self, state):
    """Return the next move based on the current state.

    Args:
        state: Current game board state as numpy array

    Returns:
        Move as (row, col) tuple
    """
    game = TicTacToe()
    game.board = state.copy()
    game.current_player = self.player_number

    # Get all legal moves
    legal_moves = game.get_legal_moves()
    if not legal_moves:
        return None

    # Check for winning move
    for move in legal_moves:
        test_game = TicTacToe()
        test_game.board = state.copy()
        test_game.current_player = self.player_number
        test_game.make_move(move)
        if test_game.is_game_over() and test_game.get_winner() == self.player_number:
            return move

    # Check for blocking opponent's winning move
    opponent = 3 - self.player_number
    for move in legal_moves:
        test_board = state.copy()
        test_board[move[0], move[1]] = opponent

        # Check if this move would make opponent win
        test_game = TicTacToe()
        test_game.board = test_board
        test_game._check_game_over()
        if test_game.is_game_over() and test_game.get_winner() == opponent:
            return move

    # If middle square is available, take it (strategic advantage)
    if (1, 1) in legal_moves:
        return (1, 1)

    # Prefer corners over sides (strategic advantage)
    corners = [(0, 0), (0, 2), (2, 0), (2, 2)]
    available_corners = [move for move in corners if move in legal_moves]
    if available_corners:
        return random.choice(available_corners)

    # Otherwise, play randomly
    return random.choice(legal_moves)

if __name__ == "__main__":
    # Example of using the default opponent
    from games.tic_tac_toe import TicTacToeUI, PlayerType, GameMode

    # Create the game UI
    game_ui = TicTacToeUI()

    # Create the default opponent and set as player 2
    default_opponent = DefaultOpponentTTT(player_number=2)
    game_ui.set_player2_agent(default_opponent)

    # Set the game mode to Human vs Semi-Intelligent
    game_ui.game_mode = GameMode.HUMAN_VS_SEMI
    game_ui.player1_type = PlayerType.HUMAN
    game_ui.player2_type = PlayerType.SEMI_INTELLIGENT

    # Run the game
    game_ui.run()

```

D.2 Default Opponent Connect 4

```

import random
import numpy as np

```

```

from games.connect4 import Connect4

class DefaultOpponentC4:
    """A semi-intelligent agent for Connect 4 that:
    - Plays a winning move if available
    - Blocks opponent's winning move if possible
    - Prefers center columns (strategic advantage)
    - Otherwise plays randomly
    """
    def __init__(self, player_number=2):
        """Initialize the agent.

        Args:
            player_number: 1 for Red, 2 for Yellow (default is 2)
        """
        self.player_number = player_number

    def get_move(self, state):
        """Return the next move based on the current state.

        Args:
            state: Current game board state as numpy array

        Returns:
            Move as column index (0-6)
        """
        game = Connect4()
        game.board = state.copy()
        game.current_player = self.player_number

        # Get all legal moves
        legal_moves = game.get_legal_moves()
        if not legal_moves:
            return None

        # Check for winning move
        for col in legal_moves:
            test_game = Connect4()
            test_game.board = state.copy()
            test_game.current_player = self.player_number
            test_game.make_move(col)
            if test_game.is_game_over() and test_game.get_winner() == self.player_number:
                return col

        # Check for blocking opponent's winning move
        opponent = 3 - self.player_number
        for col in legal_moves:
            test_game = Connect4()
            test_game.board = state.copy()
            test_game.current_player = opponent
            test_game.make_move(col)
            if test_game.is_game_over() and test_game.get_winner() == opponent:
                return col

        # Prefer middle columns
        # The closer to the middle, the higher the probability of being chosen
        weights = []
        middle = 3 # For a 7-column board, the middle is index 3
        for col in legal_moves:
            # Calculate weight based on distance from middle
            distance = abs(col - middle)
            weight = 7 - distance # Higher weight for columns closer to middle
            weights.append(weight)

        # Normalize weights to probabilities
        total_weight = sum(weights)
        probabilities = [w / total_weight for w in weights]

        # Choose column based on weights
        return random.choices(legal_moves, weights=probabilities, k=1)[0]

if __name__ == "__main__":
    # Example of using the default opponent
    from games.connect4 import Connect4UI, PlayerType, GameMode

    # Create the game UI
    game_ui = Connect4UI()

    # Create the default opponent and set as player 2
    default_opponent = DefaultOpponentC4(player_number=2)
    game_ui.set_player2_agent(default_opponent)

    # Set the game mode to Human vs Semi-Intelligent
    game_ui.game_mode = GameMode.HUMAN_VS_SEMI
    game_ui.player1_type = PlayerType.HUMAN
    game_ui.player2_type = PlayerType.SEMI_INTELLIGENT

    # Run the game
    game_ui.run()

```